

(Hopefully) Efficient Numerical Estimation of Optical Flow

Thorbjørn Djupvik, Rudolfs Sietinsons & Ivan Nygaard

Problem Statement (part I)

$$I : \Omega \times \mathbb{R} \rightarrow \mathbb{R} \quad \longleftarrow \quad \text{Video (Intensity)}$$

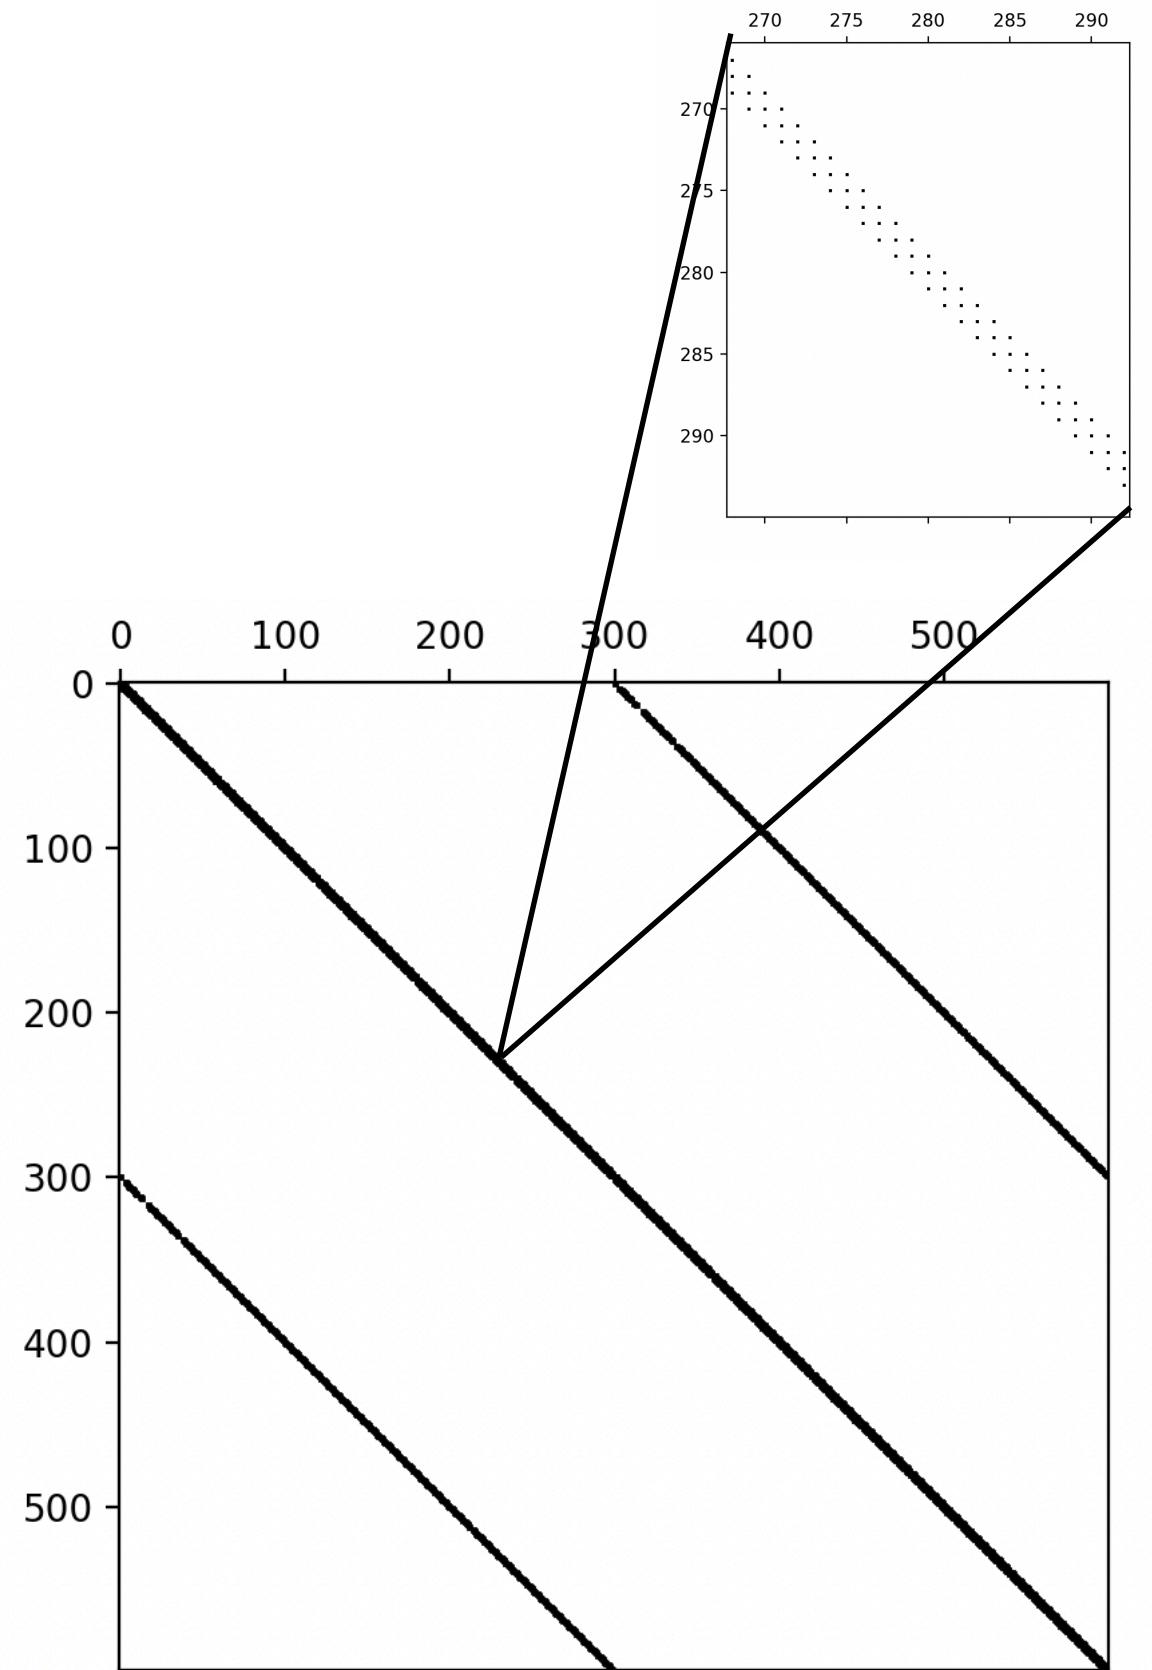
$$\mathbf{w}(x, y) := [u(x, y), v(x, y)] \quad \longleftarrow \quad \text{Optical flow-field (what we want)}$$

$$I(x(t_0 + \Delta t), y(t_0 + \Delta t), t_0 + \Delta t) = I(x(t_0), y(t_0), t_0) \quad \longleftarrow \quad \text{Brightness consistency}$$

$$u\partial_x I + v\partial_y I = -\partial_t I \quad \longleftarrow \quad \text{PDE}$$

$$\left. \begin{aligned} & \left[(\partial_x I)_{ij}^2 + \frac{4\lambda}{h^2} \right] u_{ij} - \frac{\lambda}{h^2} \left[u_{i+1,j} + u_{i-1,j} + u_{i,j+1} + u_{i,j-1} \right] + (\partial_y I)_{ij} (\partial_x I)_{ij} v_{ij} = -(\partial_t I)_{ij} (\partial_x I)_{ij}, \\ & \left[(\partial_y I)_{ij}^2 + \frac{4\lambda}{h^2} \right] v_{ij} - \frac{\lambda}{h^2} \left[v_{i+1,j} + v_{i-1,j} + v_{i,j+1} + v_{i,j-1} \right] + (\partial_y I)_{ij} (\partial_x I)_{ij} u_{ij} = -(\partial_t I)_{ij} (\partial_y I)_{ij}. \end{aligned} \right\} \quad \longleftarrow \quad \text{Discretized equations}$$

Problem Statement (part II)

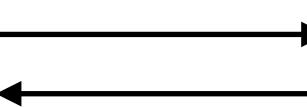


$(2mn) \times (2mn)$

$(2mn)$

$(2mn)$

$$\begin{array}{ll} u_{11} & \text{rhsu}_{11} \\ u_{12} & \text{rhsu}_{12} \\ u_{13} & \text{rhsu}_{13} \\ \dots & \dots \\ v_{11} & \text{rhsv}_{11} \\ v_{12} & \text{rhsv}_{12} \\ v_{13} & \text{rhsv}_{13} \\ \dots & \dots \end{array} = \begin{array}{l} \dots \end{array}$$

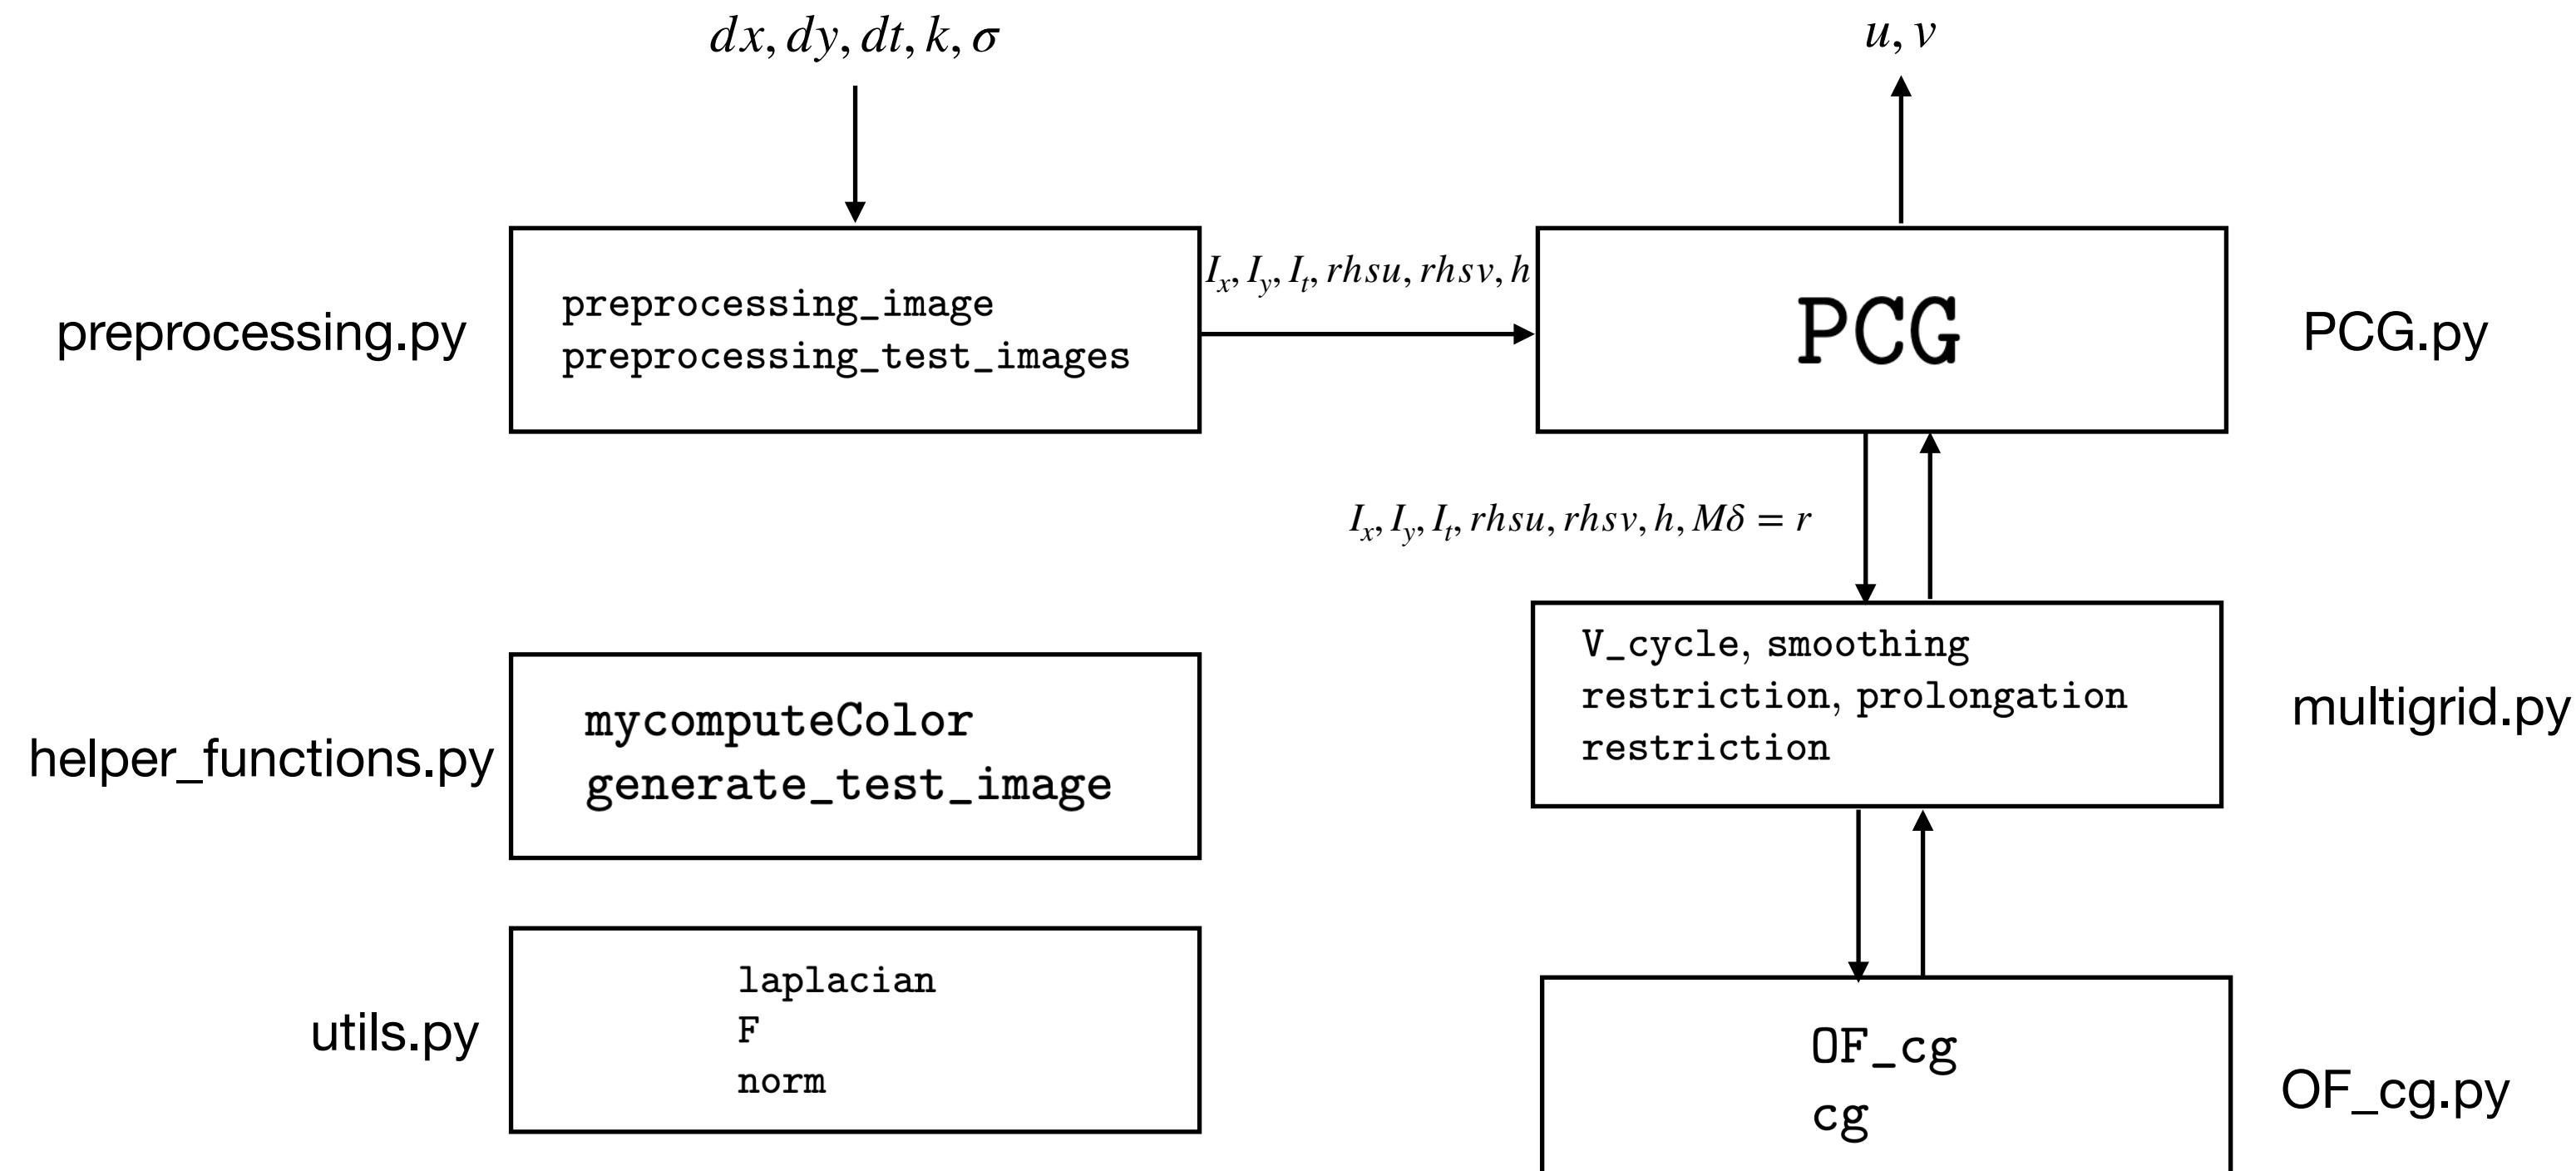


$$\underbrace{\begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix}}_A \underbrace{\begin{bmatrix} \mathbf{u} \\ \mathbf{v} \end{bmatrix}}_x = \underbrace{\begin{bmatrix} \mathbf{b}_1 \\ \mathbf{b}_2 \end{bmatrix}}_b$$

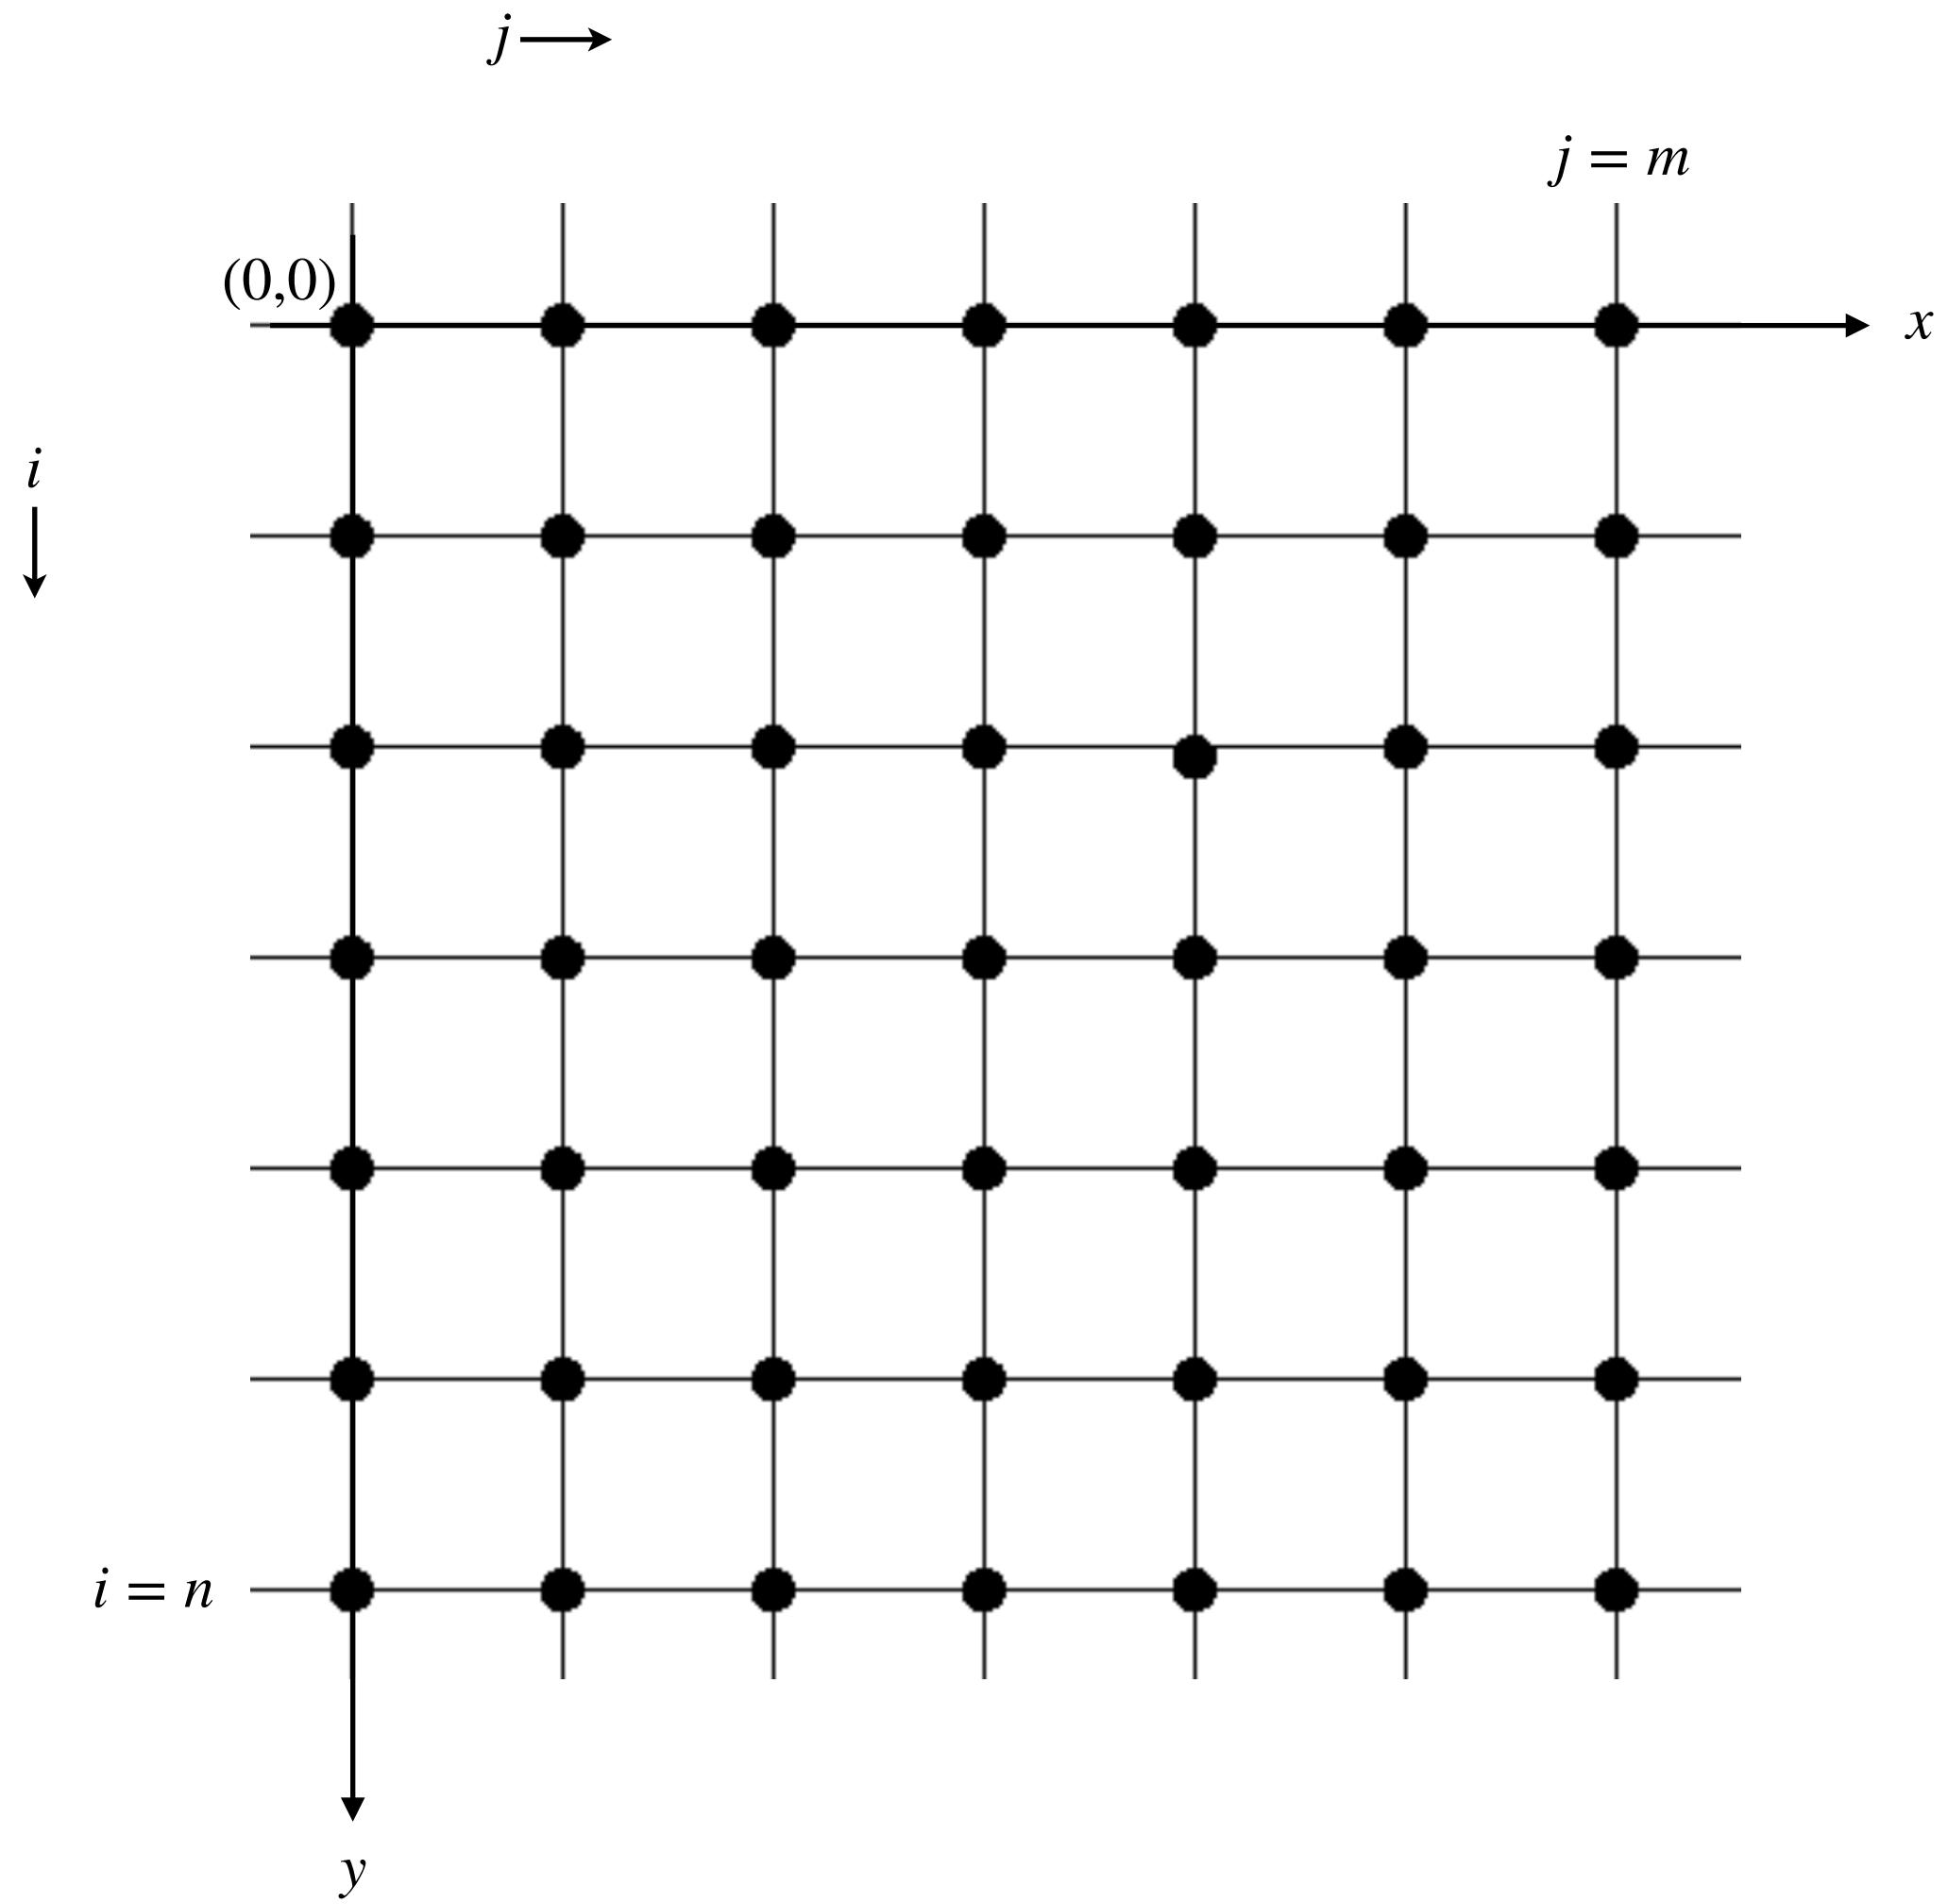
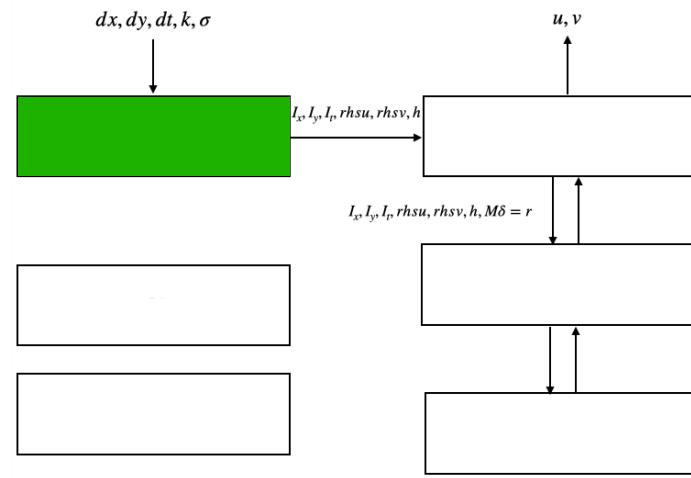
$$\left. \begin{aligned} & \left[(\partial_x I)_{ij}^2 + \frac{4\lambda}{h^2} \right] u_{ij} - \frac{\lambda}{h^2} [u_{i+1,j} + u_{i-1,j} + u_{i,j+1} + u_{i,j-1}] + (\partial_y I)_{ij} (\partial_x I)_{ij} v_{ij} = -(\partial_t I)_{ij} (\partial_x I)_{ij}, \\ & \left[(\partial_y I)_{ij}^2 + \frac{4\lambda}{h^2} \right] v_{ij} - \frac{\lambda}{h^2} [v_{i+1,j} + v_{i-1,j} + v_{i,j+1} + v_{i,j-1}] + (\partial_y I)_{ij} (\partial_x I)_{ij} u_{ij} = -(\partial_t I)_{ij} (\partial_y I)_{ij}. \end{aligned} \right\}$$

The matrix A can be proven to be SPD, this has some important consequences for the implementation.

Workflow



preprocessing.py



$$\partial_x I_k(x_i, y_j) := \begin{cases} I_k(x_{i+1}, y_j) - I_k(x_i, y_j), & \text{if } i < m, \\ I_k(x_m, y_j) - I_k(x_{m-1}, y_j), & \text{if } i = m. \end{cases}$$

$|I0[:, :-1] = (|I0[:, 1:] - |I0[:, :-1]) / dx$

← Frwrd. Diff.

$$\partial_x I_k(x_i, y_j) := \begin{cases} I_k(x_i, y_{j+1}) - I_k(x_i, y_j), & \text{if } j < n, \\ I_k(x_i, y_n) - I_k(x_i, y_{n-1}), & \text{if } j = n. \end{cases}$$

$|I0[:, -1] = (|I0[:, -1] - |I0[:, -2]) / dx$

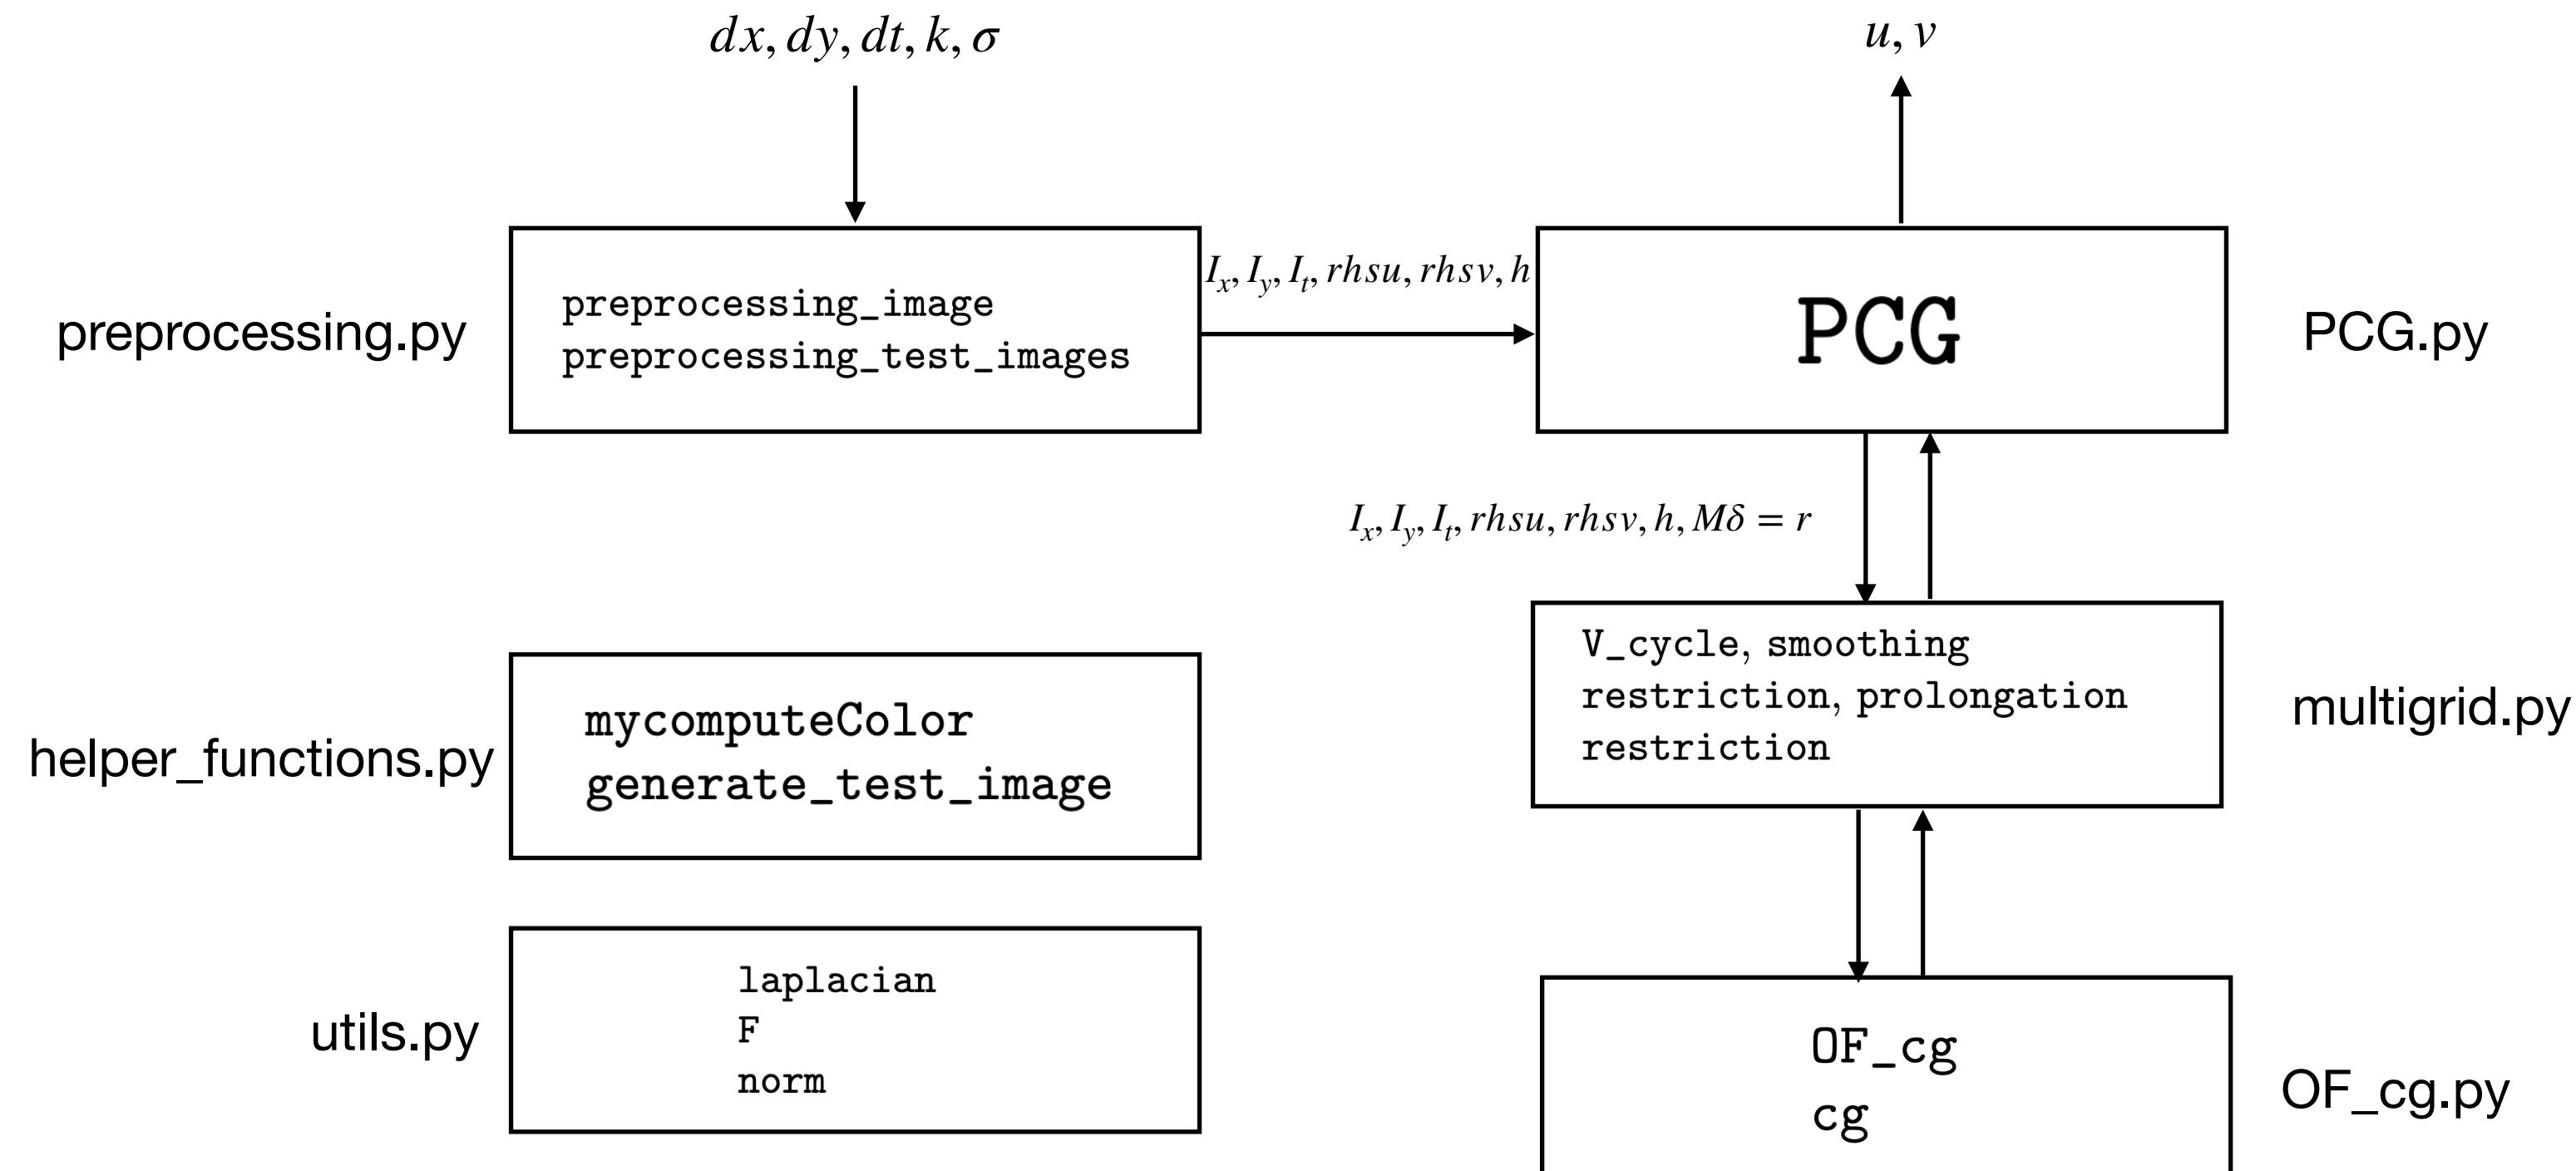
← Bcwr. Diff.

$$\partial_t I(x_i, y_j) := I_1(x_i, y_j) - I_0(x_i, y_j)$$

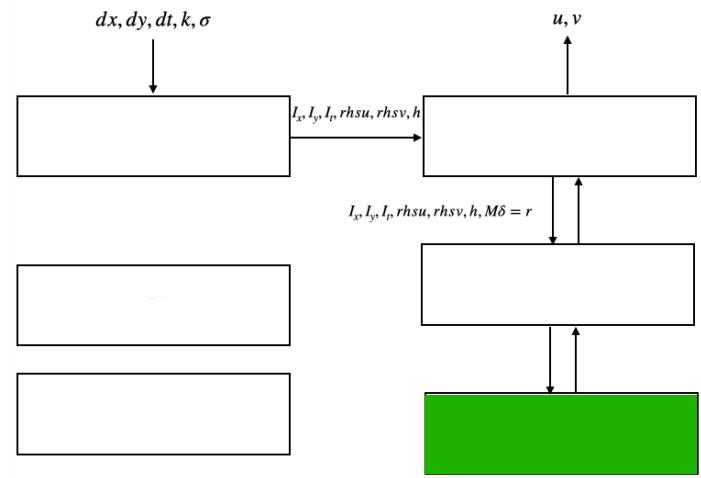
$|It = (|I1 - |I0|) / dt$

← Frwrd. Diff
«Frame difference»

Workflow



OF_cg.py



Algorithm 1 CG

- 1: Compute $r_0 := b - Ax_0$, $p_0 := r_0$
 - 2: **for** $j = 0, 1, \dots$, until convergence **do**
 - 3: $\alpha_j \leftarrow \frac{(r_j, r_j)}{(Ap_j, p_j)}$
 - 4: $x_{j+1} \leftarrow x_j + \alpha_j p_j$
 - 5: $r_{j+1} \leftarrow r_j - \alpha_j Ap_j$
 - 6: $\beta_j \leftarrow \frac{(r_{j+1}, r_{j+1})}{(r_j, r_j)}$
 - 7: $p_{j+1} \leftarrow r_{j+1} + \beta_j p_j$
 - 8: **end for**
-

Two implementations, one sparse and one that works directly on the grid.

```

A_11 = sp.diags(Ix.ravel() ** 2) + L
A_22 = sp.diags(Iy.ravel() ** 2) + L
A_12 = sp.diags(Ixy.ravel())
A_21 = A_12.copy()

A = sp.block_array([[A_11, A_12], [A_21, A_22]]).tocsr()
  
```

← Sparse CG

```

# Residuals (For experiments)
relative_residuals_arr = np.zeros(maxitr)

# Initialize
Fu, Fv = F(u0, v0, Ix, Iy, lam, h)

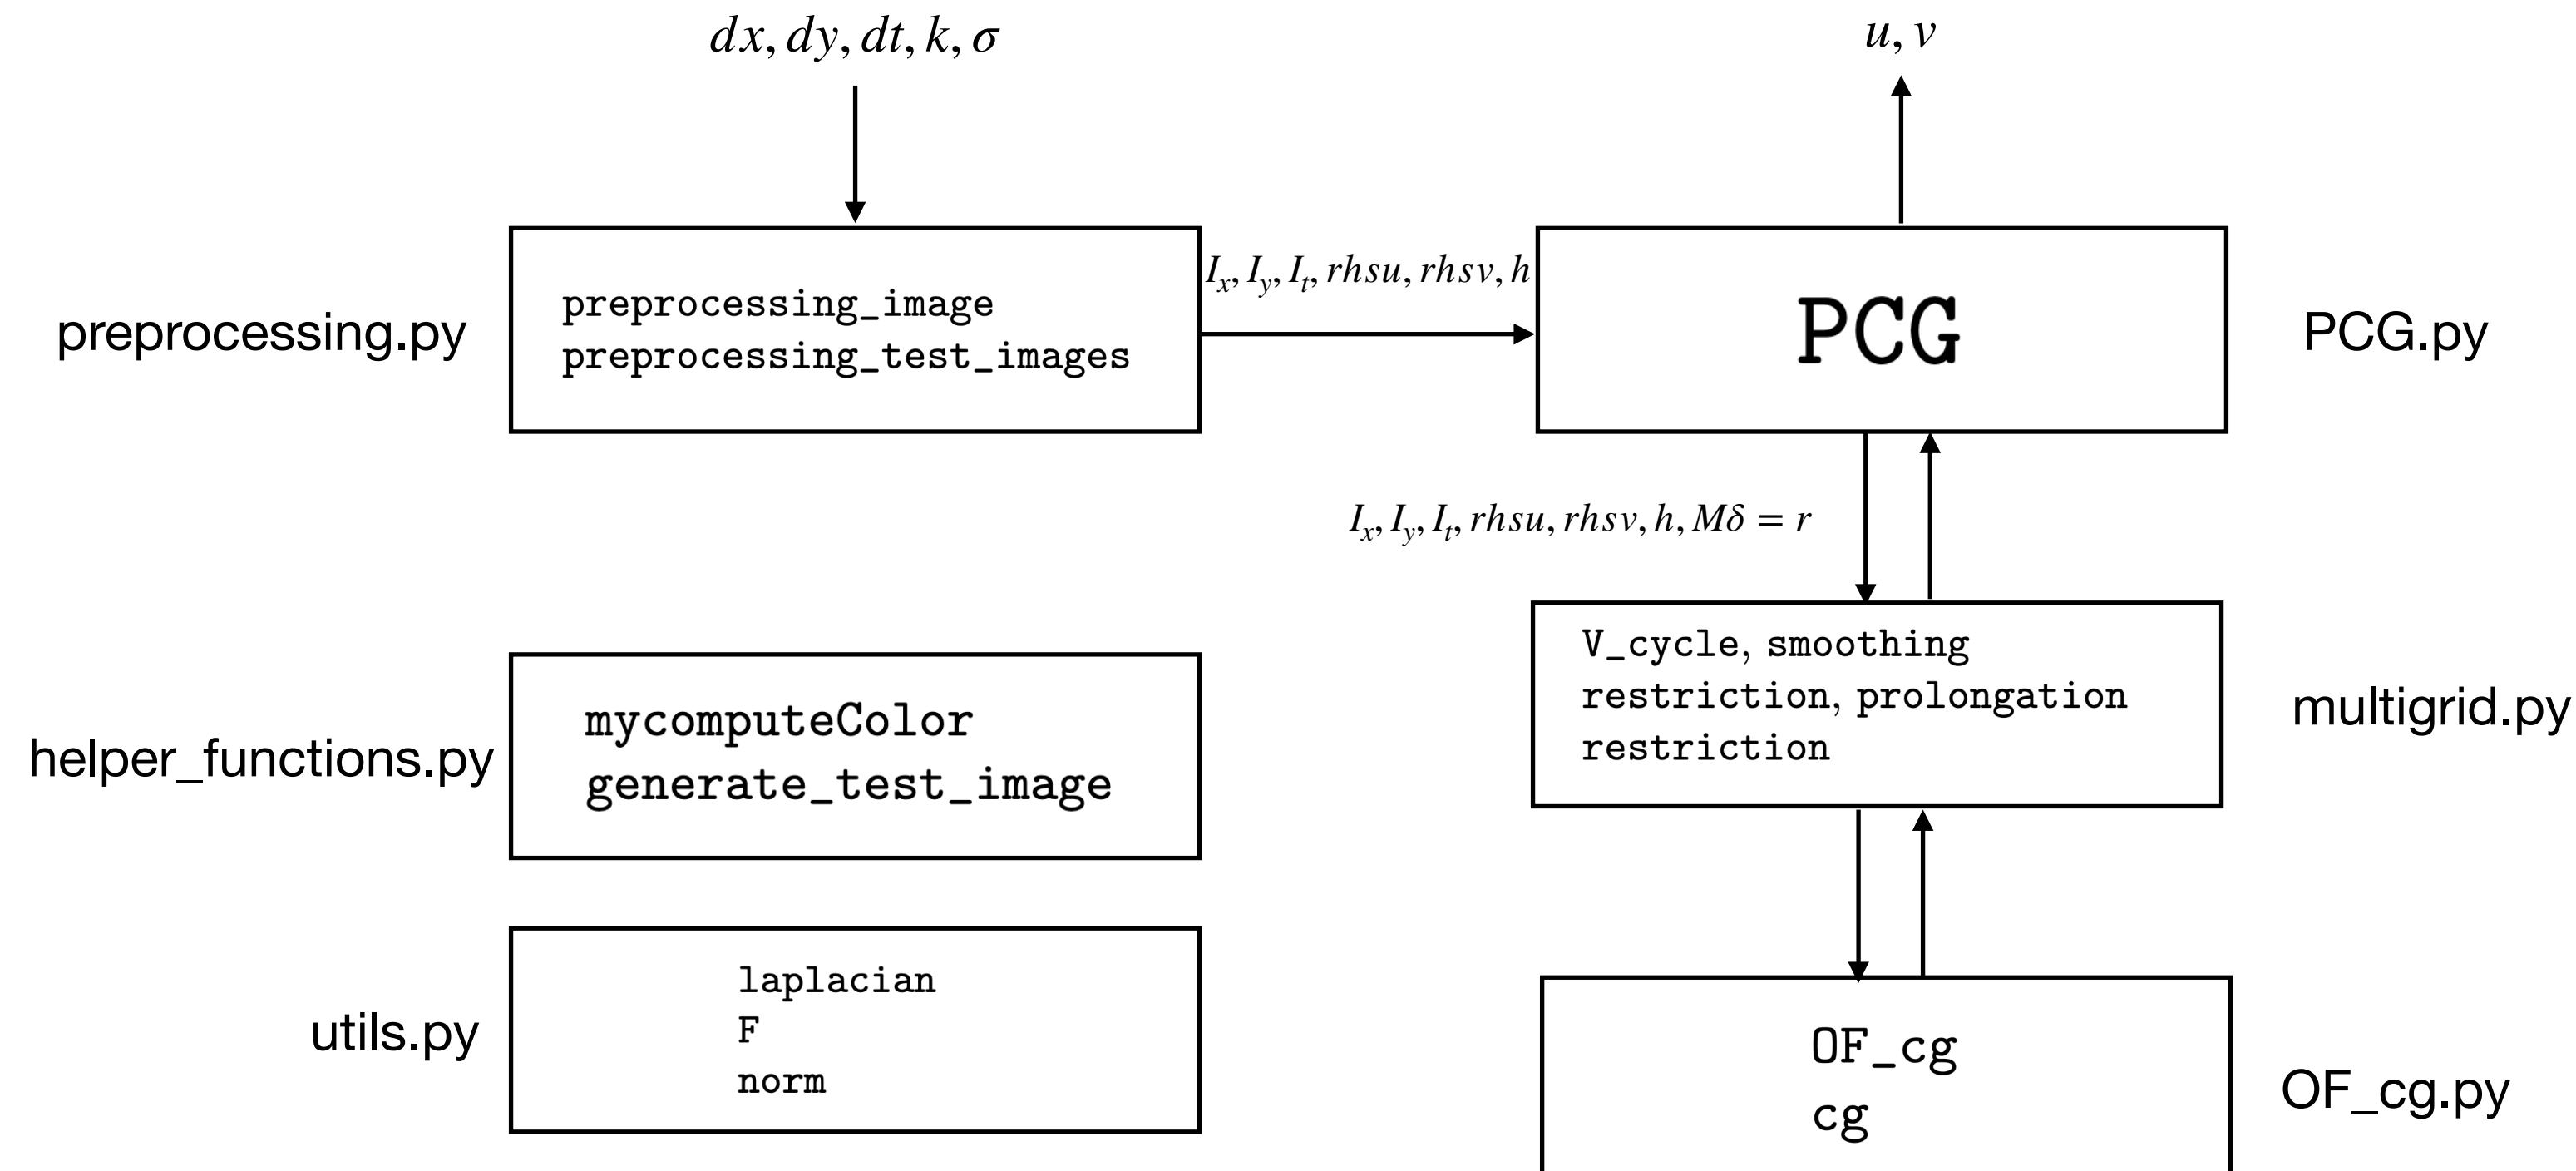
# r
ru = np.copy(rhs_u - Fu)
rv = np.copy(rhs_v - Fv)
# x
u = np.copy(u0)
v = np.copy(v0)
# p
pu = np.copy(ru)
pv = np.copy(rv)

# Calculate the norm of r
rr0 = norm(ru, rv) ** 2
rk1_rk1 = rr0

def F(
    u: np.ndarray,
    v: np.ndarray,
    Ix: np.ndarray,
    Iy: np.ndarray,
    lam: float,
    h: float,
) -> tuple[np.ndarray, np.ndarray]:
    Ixy = Ix*Iy
    Fu = Ix**2 * u + Ixy * v - lam * laplacian(u, h)
    Fv = Iy**2 * v + Ixy * u - lam * laplacian(v, h)
    return Fu, Fv
  
```

← Grid CG

Workflow



Multigrid V-cycle (part I)

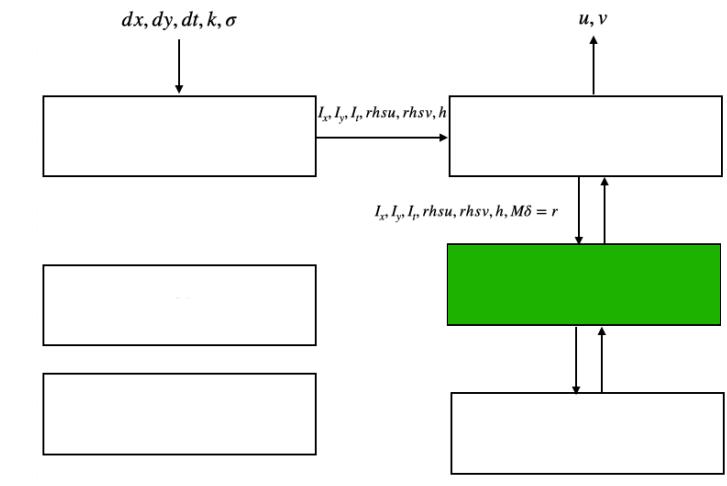
```
def V_cycle(
    u0: np.ndarray,
    v0: np.ndarray,
    Ix: np.ndarray,
    Iy: np.ndarray,
    lam: float,
    rhs_u: np.ndarray,
    rhs_v: np.ndarray,
    s1: int,
    s2: int,
    level: int,
    max_level: int,
) -> tuple[np.ndarray, np.ndarray]:
    """
    V-cycle for the optical flow problem

    Args:
    ---
    u0 : np.ndarray
        Initial guess for u
    v0 : np.ndarray
        Initial guess for v
    Ix : np.ndarray
        x-derivative of the first frame
    Iy : np.ndarray
        y-derivative of the first frame
    lam : float
        Penalty term
    rhs_u : np.ndarray
        RHS in eq for u
    rhs_v : np.ndarray
        RHS in eq for v
    s1 : int
        Number of pre-smoothings
    s2 : int
        Number of post-smoothings
    level : int
        Current level
    max_level : int
        Total number of levels

    Returns:
    ---
    tuple[np.ndarray, np.ndarray, list[float]]
        Numerical solution for u, v
    """

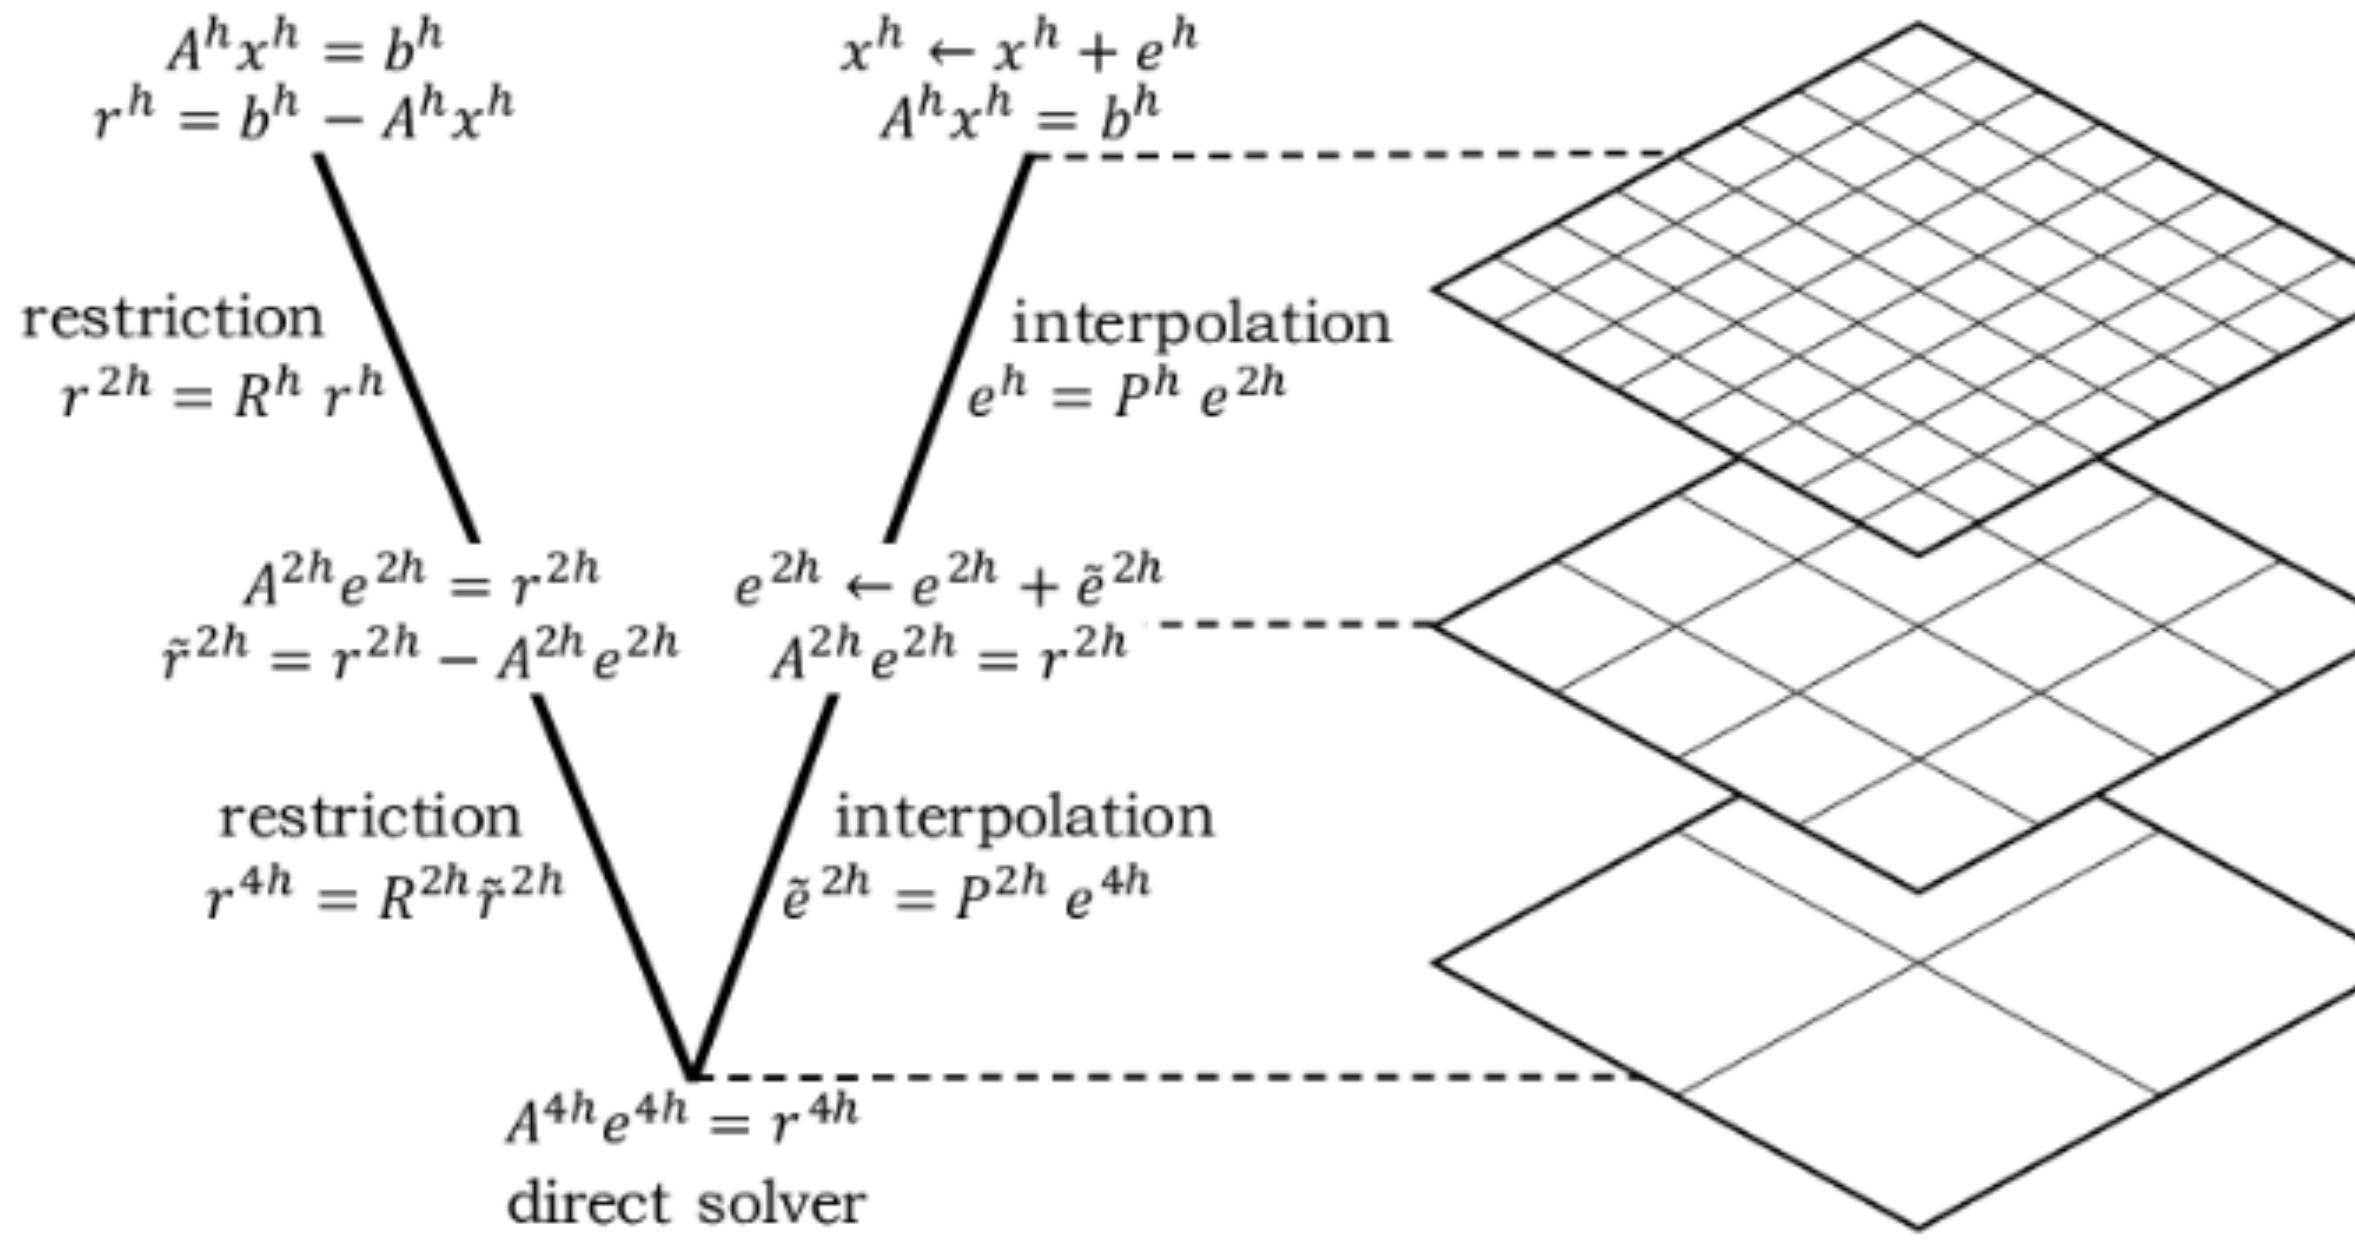
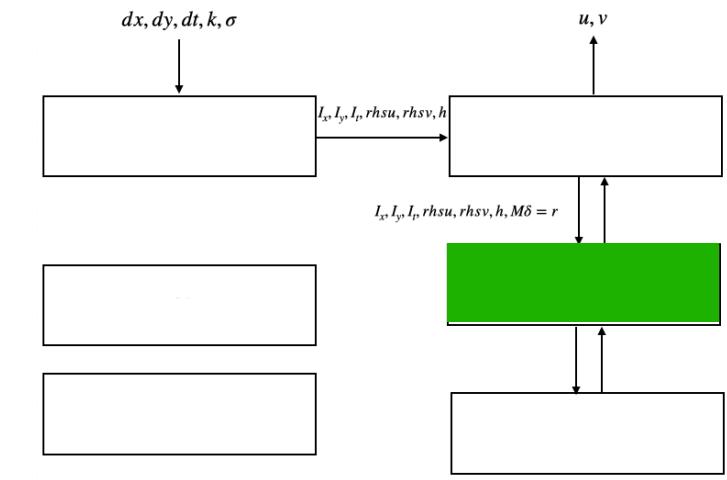
```

← V-cycle routine



The method is implemented recursively and calls the sub-routines red_update, black_update, smoothing, prolongation, restriction and interpolation.

Multigrid V-cycle (part II)



```

assert ru_h.shape == rv_h.shape == Ix.shape == Iy.shape
n, m = ru_h.shape

ru_2h = (
    ru_h[0 : n - 1 : 2, 0 : m - 1 : 2] # upper left
    + ru_h[1:n:2, 0 : m - 1 : 2] # lower left
    + ru_h[0 : n - 1 : 2, 1:m:2] # upper right
    + ru_h[1:n:2, 1:m:2] # lower right
) / 4
rv_2h = (
    rv_h[0 : n - 1 : 2, 0 : m - 1 : 2] # upper left
    + rv_h[1:n:2, 0 : m - 1 : 2] # lower left
    + rv_h[0 : n - 1 : 2, 1:m:2] # upper right
    + rv_h[1:n:2, 1:m:2] # lower right
) / 4
Ix2h = (
    Ix[0 : n - 1 : 2, 0 : m - 1 : 2] # upper left
    + Ix[1:n:2, 0 : m - 1 : 2] # lower left
    + Ix[0 : n - 1 : 2, 1:m:2] # upper right
    + Ix[1:n:2, 1:m:2] # lower right
) / 4
Iy2h = (
    Iy[0 : n - 1 : 2, 0 : m - 1 : 2] # upper left
    + Iy[1:n:2, 0 : m - 1 : 2] # lower left
    + Iy[0 : n - 1 : 2, 1:m:2] # upper right
    + Iy[1:n:2, 1:m:2] # lower right
) / 4

```

Restriction

```

n, m = e_2h.shape

e_h = np.zeros((2 * n, 2 * m))
k, d = e_h.shape

# Upper left
e_h[0 : k - 1 : 2, 0 : d - 1 : 2] = e_2h
# Lower left
e_h[1:k:2, 0 : d - 1 : 2] = e_2h
# Upper right
e_h[0 : k - 1 : 2, 1:d:2] = e_2h
# Lower right
e_h[1:k:2, 1:d:2] = e_2h

return e_h

```

Prolongation

Algorithm 3 Red-Black

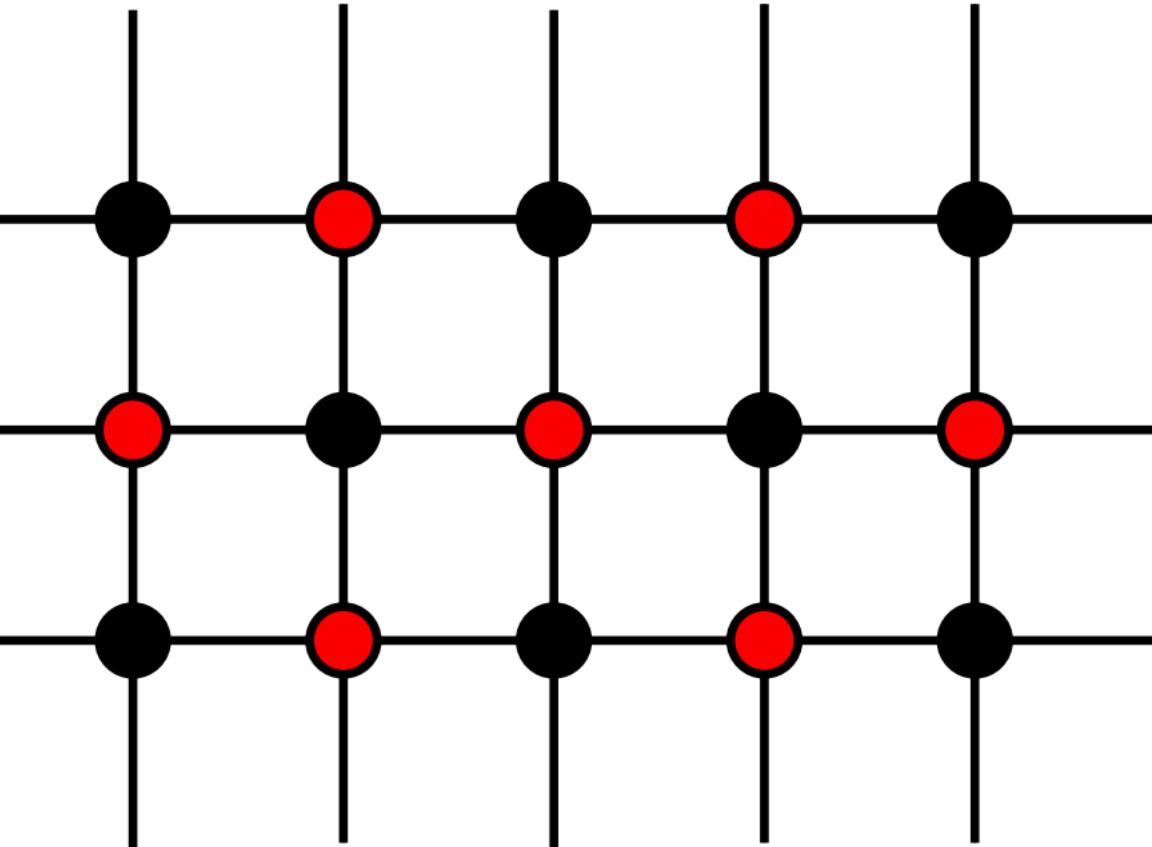
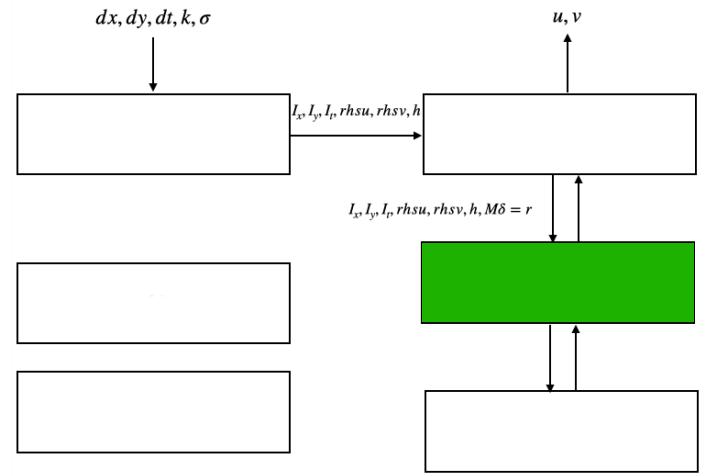
- 1: **Red-Update(u)**
- 2: **Black-Update(v)**
- 3: **Red-Update(v)**
- 4: **Black-Update(u)**

```

241 def red_update(
242     w: np.ndarray,
243     p: np.ndarray,
244     Iw: np.ndarray,
245     Ip: np.ndarray,
246     lam: float,
247     rhs: np.ndarray,
248     h: float,
249 ) -> None:
250     """
251     Red update using Red-Black Gauss-Seidel
252     Note: This function mutates w
253
254     Args:
255     ---
256     w : np.ndarray
257         The padded array to update
258     p : np.ndarray
259         The padded other array
260     Iw : np.ndarray
261         If w=u -> Iw=Ix
262         If w=v -> Iw=Iy
263     Ip : np.ndarray
264         If w=u -> Ip=Ly
265         If w=v -> Ip=Lx
266     lam : float
267         Penalty term
268     rhs : np.ndarray
269         RHS in eq
270     h : int
271         Stepsize corresponding to the level
272
273     Returns:
274     ---
275     None
276     """
277     n, m = w.shape
278     k, d = Iw.shape
279
280     # Corner update
281     w[1 : n - 1 : 2, 1 : m - 1 : 2] = (
282         rhs[0:k:2, 0:d:2]
283         - Ip[0:k:2, 0:d:2] * Iw[0:k:2, 0:d:2] * p[1 : n - 1 : 2, 1 : m - 1 : 2]
284         + lam
285         / h**2
286         * (
287             # Left
288             w[1 : n - 1 : 2, 0 : m - 2 : 2]
289             # Right
290             + w[1 : n - 1 : 2, 2:m:2]
291             # Up
292             + w[0 : n - 2 : 2, 1 : m - 1 : 2]
293             # Down
294             + w[2:n:2, 1 : m - 1 : 2]
295         )
296     ) / (Iw[0:k:2, 0:d:2]**2 + 4 * lam / h**2)
297
298     # Interior corner update
299     w[2 : n - 1 : 2, 2 : m - 1 : 2] = (
300         rhs[1:k:2, 1:d:2]
301         - Ip[1:k:2, 1:d:2] * Iw[1:k:2, 1:d:2] * p[2 : n - 1 : 2, 2 : m - 1 : 2]
302         + lam
303         / h**2
304         * (
305             # Left
306             w[2 : n - 1 : 2, 1 : m - 2 : 2]
307             # Right
308             + w[2 : n - 1 : 2, 3:m:2]
309             # Up
310             + w[1 : n - 2 : 2, 2 : m - 1 : 2]
311             # Down
312             + w[3:m:2, 2 : m - 1 : 2]
313         )
314     ) / (Iw[1:k:2, 1:d:2]**2 + 4 * lam / h**2)

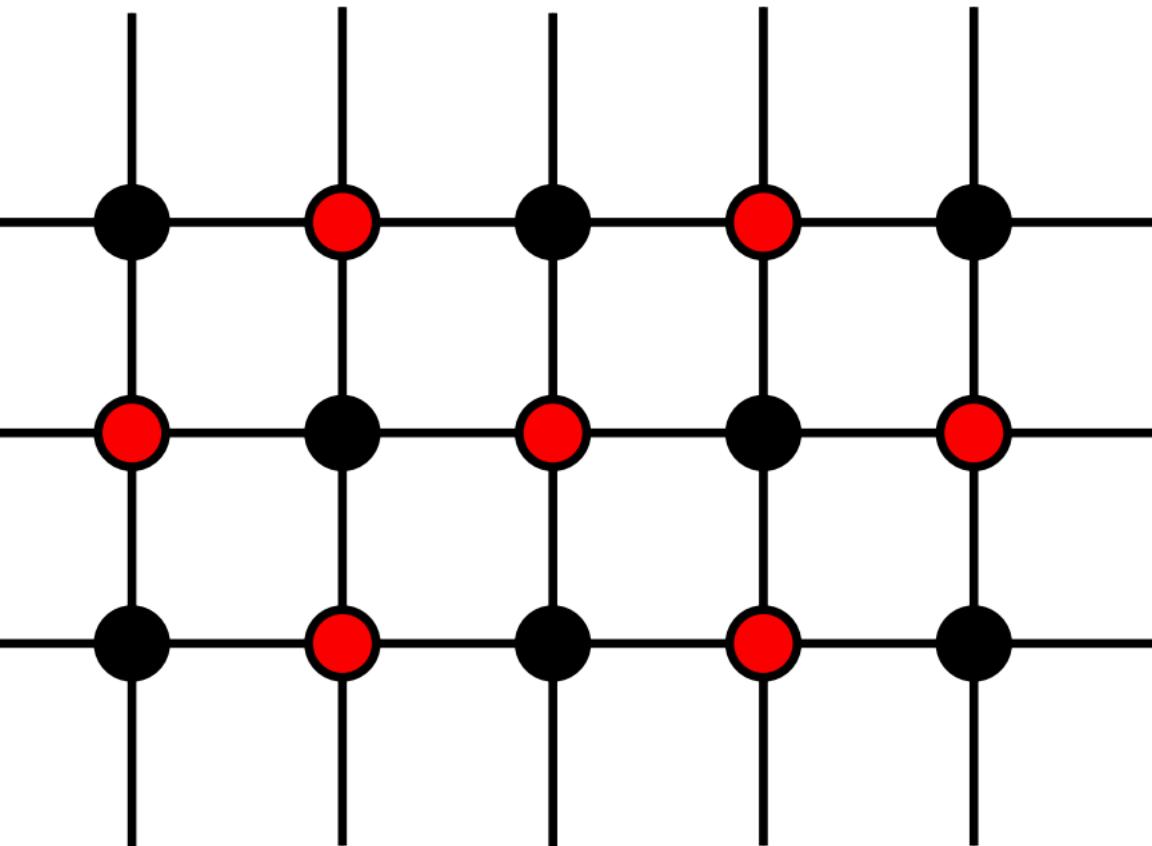
```

Multigrid V-cycle (part III)



$$u(x, y)$$

Smoother (red)

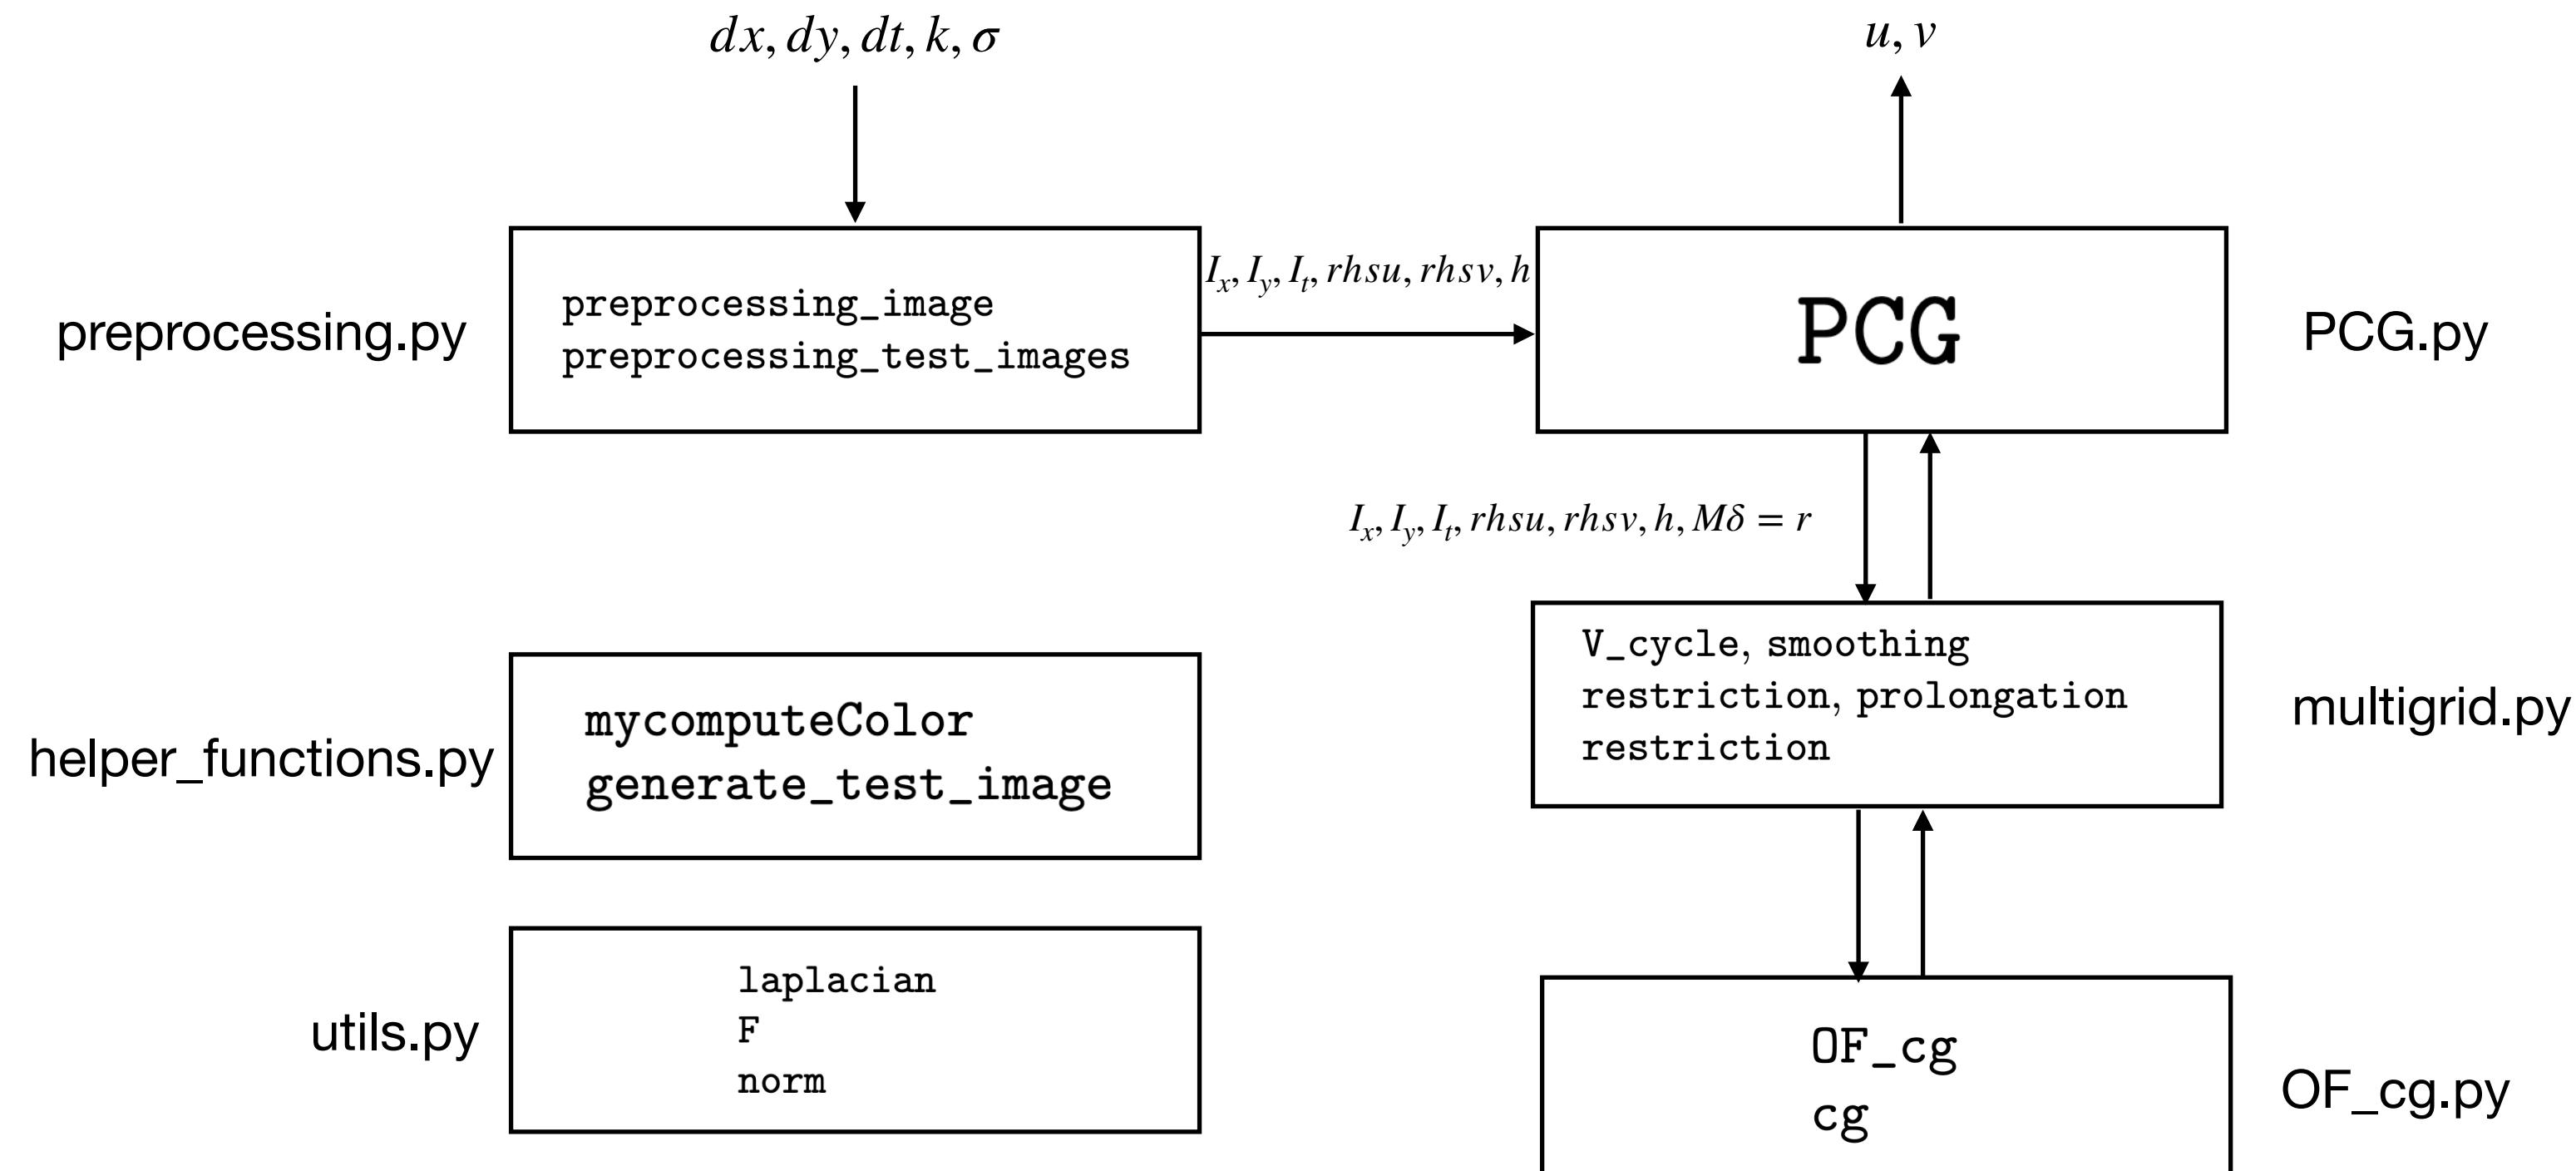


$$v(x, y)$$

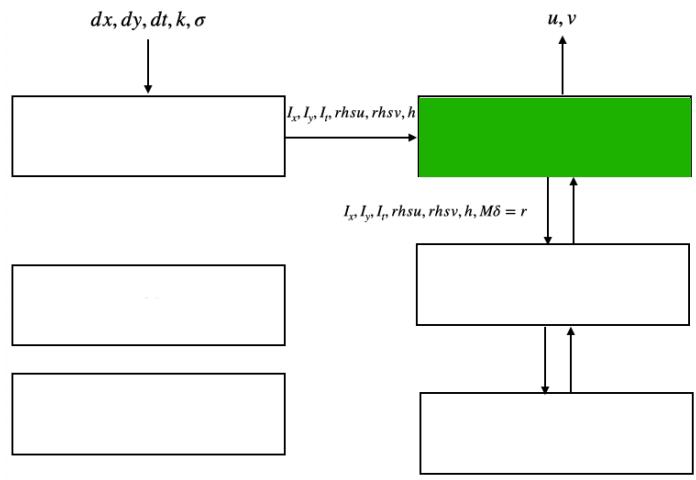
Red
Black

Red
Black

Workflow



PCG



Algorithm 1 CG

```

1: Compute  $r_0 := b - Ax_0$ ,  $p_0 := r_0$ 
2: for  $j = 0, 1, \dots$ , until convergence do
3:    $\alpha_j \leftarrow \frac{(r_j, r_j)}{(Ap_j, p_j)}$ 
4:    $x_{j+1} \leftarrow x_j + \alpha_j p_j$ 
5:    $r_{j+1} \leftarrow r_j - \alpha_j Ap_j$ 
6:    $\beta_j \leftarrow \frac{(r_{j+1}, r_{j+1})}{(r_j, r_j)}$ 
7:    $p_{j+1} \leftarrow r_{j+1} + \beta_j p_j$ 
8: end for
  
```

M needs to be symmetric positive definite!

Algorithm 2 PCG

```

1: Compute  $r_0 := b - Ax_0$ ,  $z_0 := M^{-1}r_0$ ,  $p_0 := z_0$ 
2: for  $j = 0, 1, \dots$ , until convergence do
3:    $\alpha_j \leftarrow \frac{(r_j, z_j)}{(Ap_j, p_j)}$ 
4:    $x_{j+1} \leftarrow x_j + \alpha_j p_j$ 
5:    $r_{j+1} \leftarrow r_j - \alpha_j Ap_j$ 
6:    $z_{j+1} \leftarrow M^{-1}r_{j+1}$ 
7:    $\beta_j \leftarrow \frac{(r_{j+1}, z_{j+1})}{(r_j, z_j)}$ 
8:    $p_{j+1} \leftarrow z_{j+1} + \beta_j p_j$ 
9: end for
  
```

Results (part I)

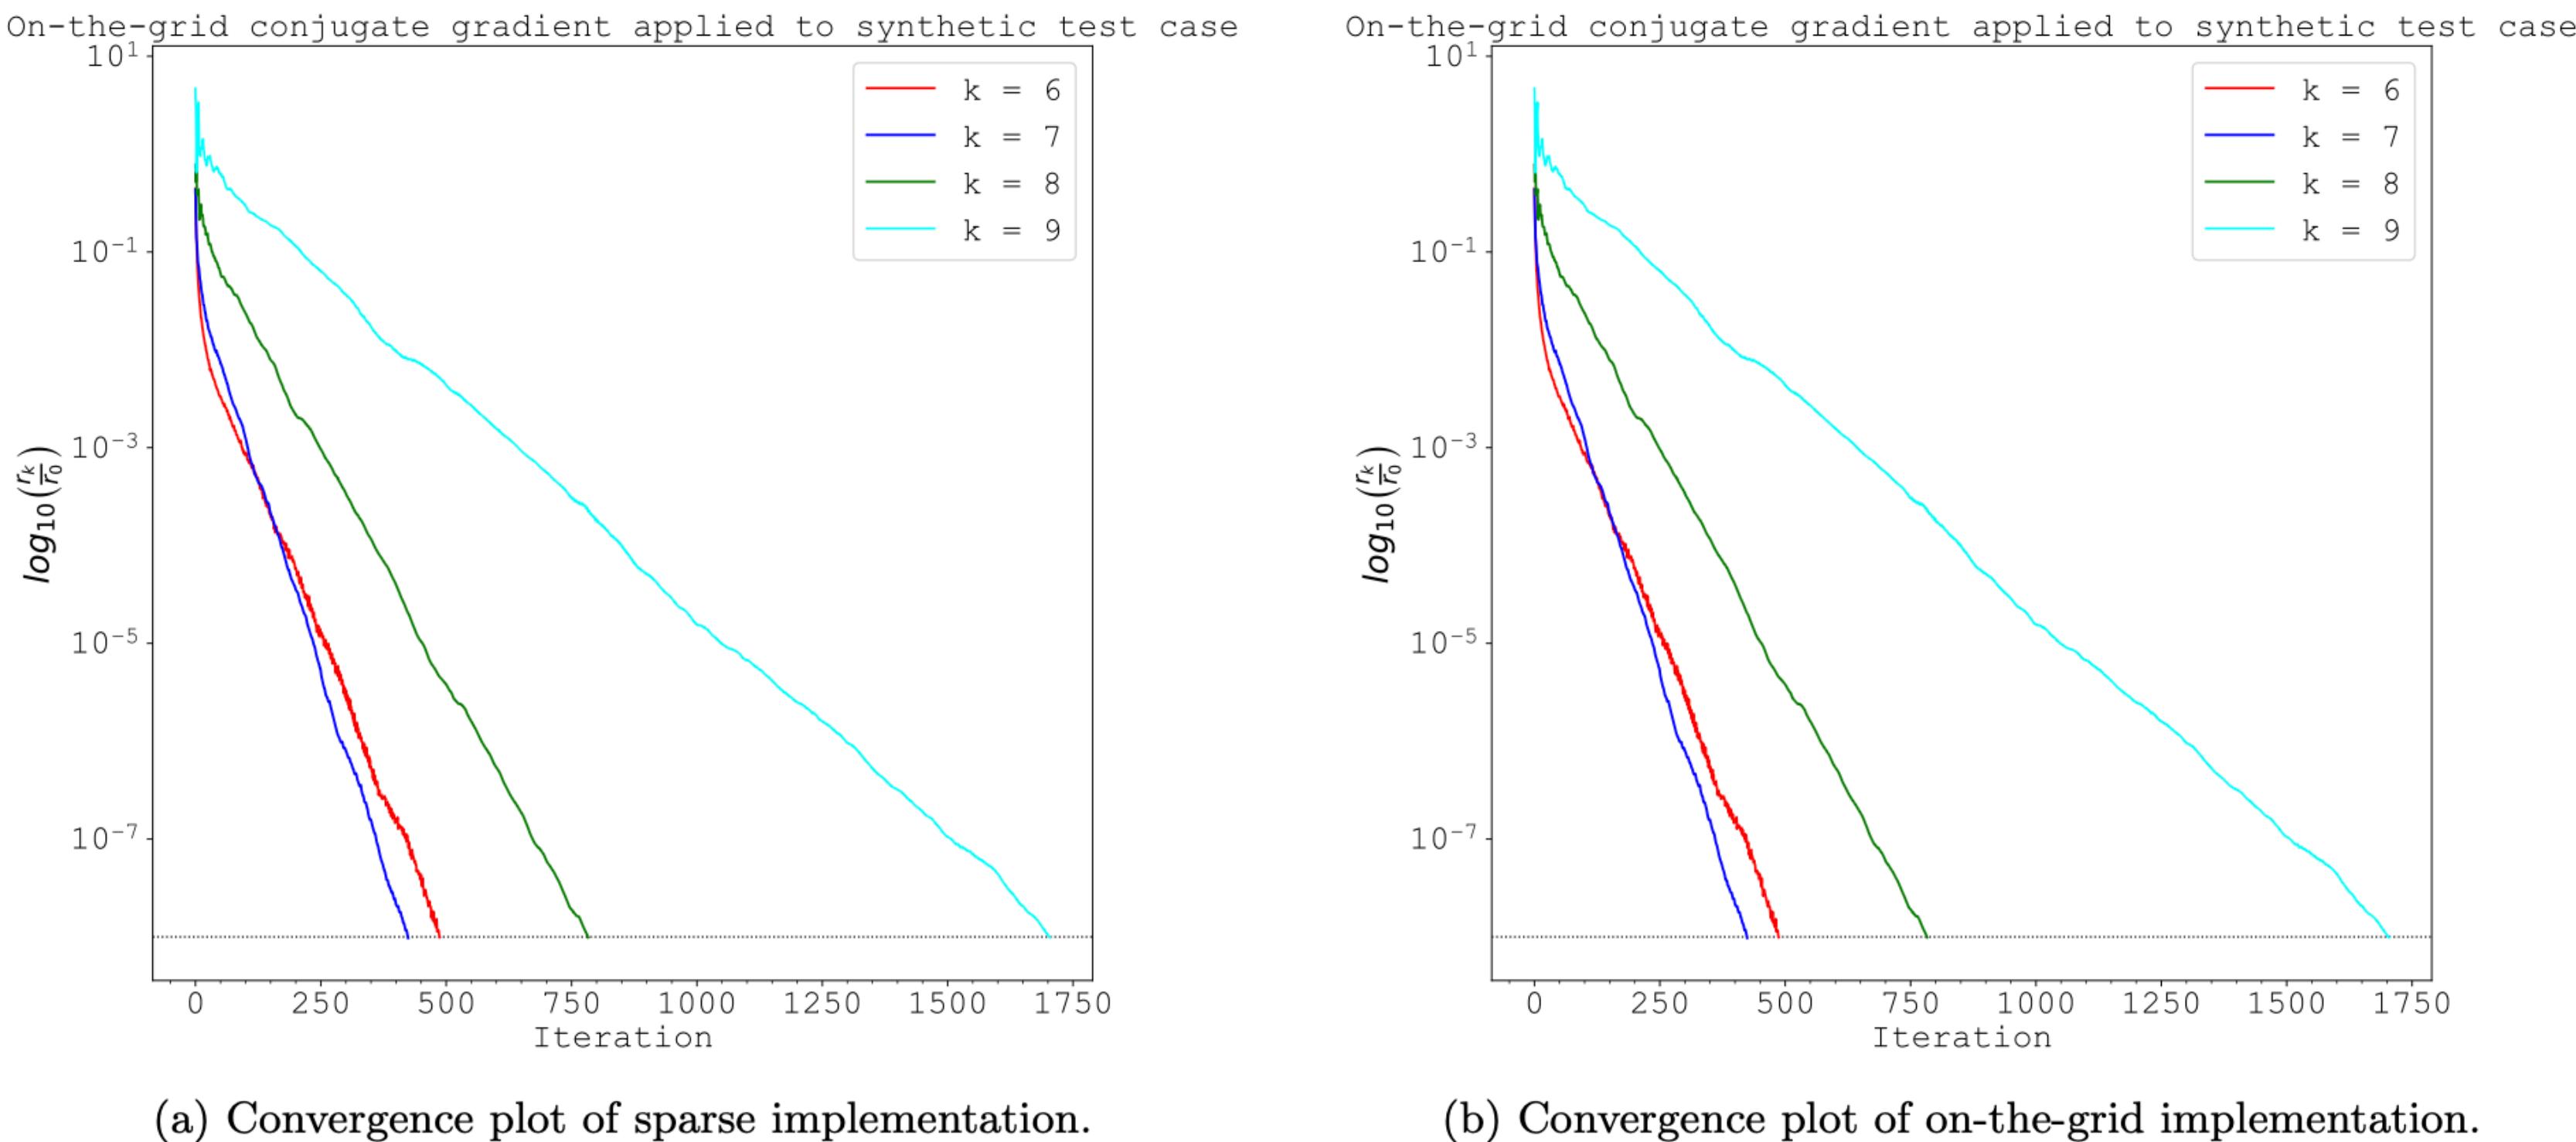


Figure 2: Convergence plots of sparse vs on-the-grid implementations of the Conjugate Gradient algorithm, without multigrid, tested on two synthetic image frames depicting two Gaussian's circling each other. The image sizes were $2^k \times 2^k$, for $k = 6, 7, 8, 9$.

k	t_{SCG} [s]	t_{GCG} [s]	t_{SCG}/t_{GCG}
6	0.047	0.086	0.546
7	0.131	0.118	1.110
8	0.988	0.815	1.212
9	9.022	9.281	0.972

Table 1: Timing comparison for sparse (t_{SCG}) vs. on-the-grid (t_{GCG}) implementations of the Conjugate Gradient algorithm without multigrid, tested on two synthetic image frames depicting two Gaussians circling each other. Image sizes are $2^k \times 2^k$ for $k = 6, 7, 8, 9$.

Results (part II)

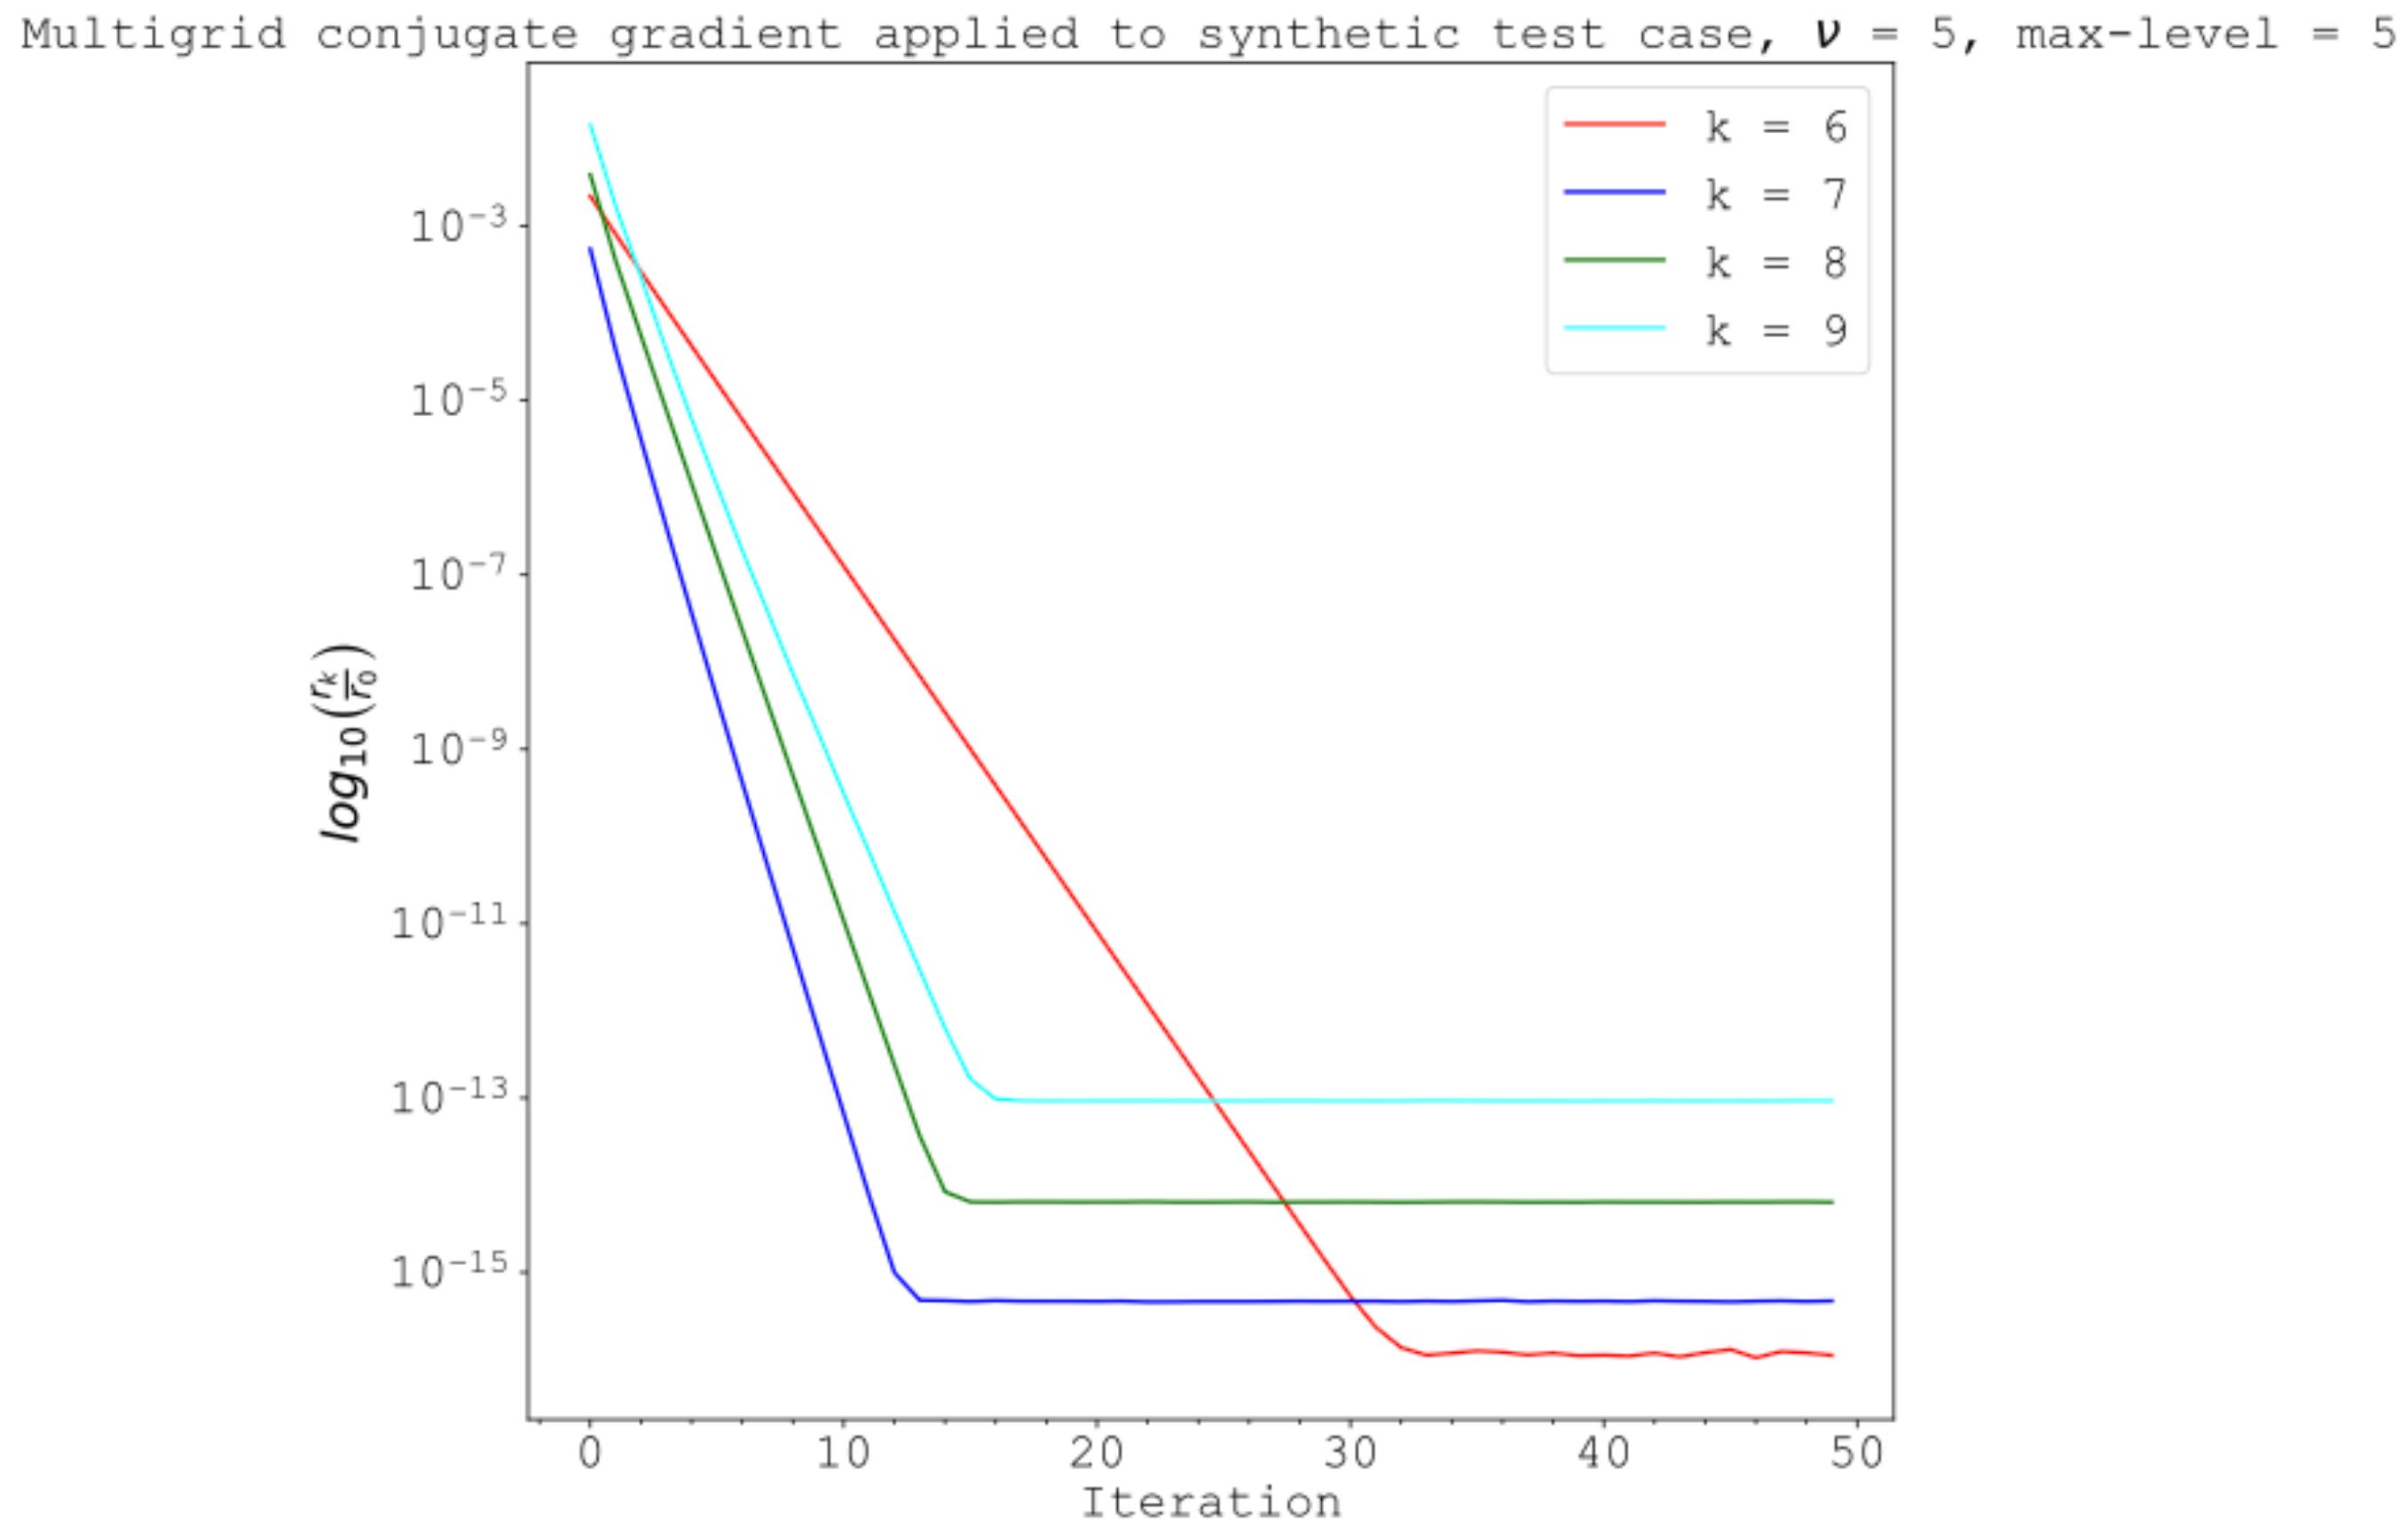
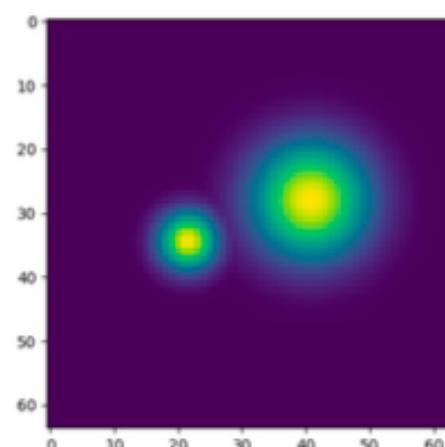
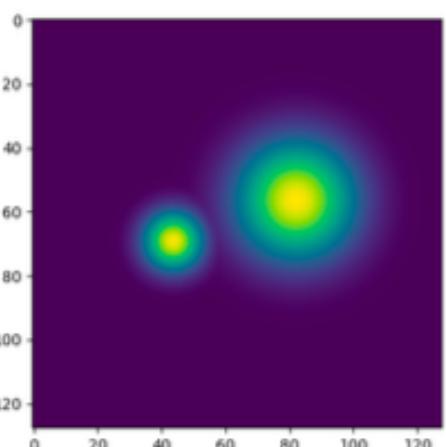


Figure 3: Convergence plot for multigrid conjugate gradient applied to the synthetic test case of two Gaussian's circling each other for image sizes of $2^k \times 2^k$ for $k = 6, 7, 8, 9$.

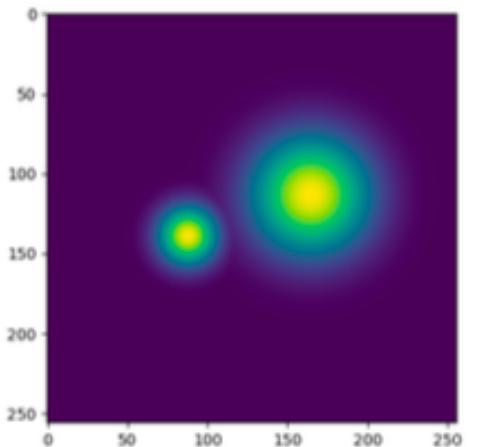
Results (part III)



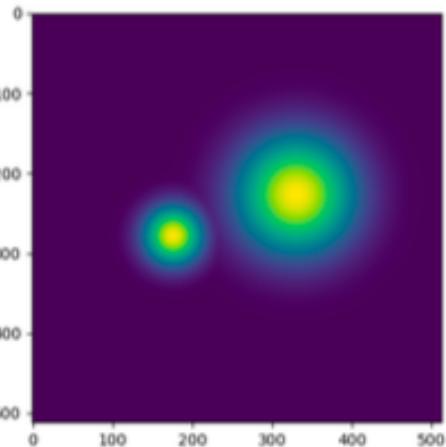
(a) $k = 6$



(b) $k = 7$

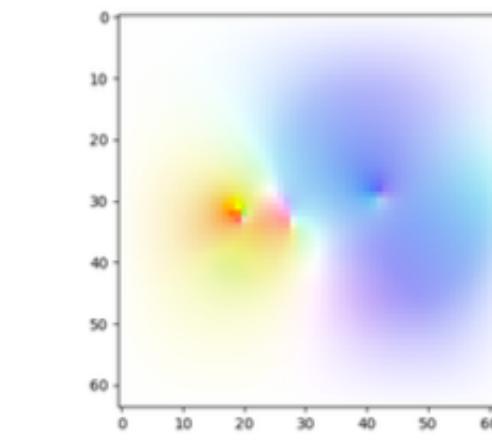


(c) $k = 8$

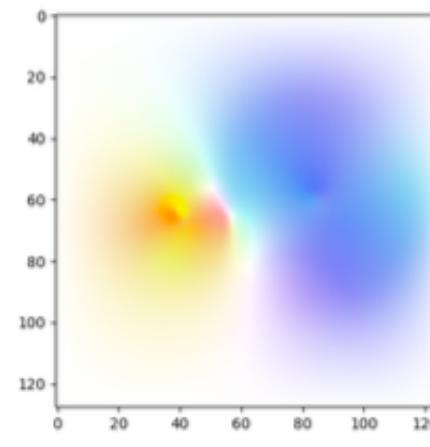


(d) $k = 9$

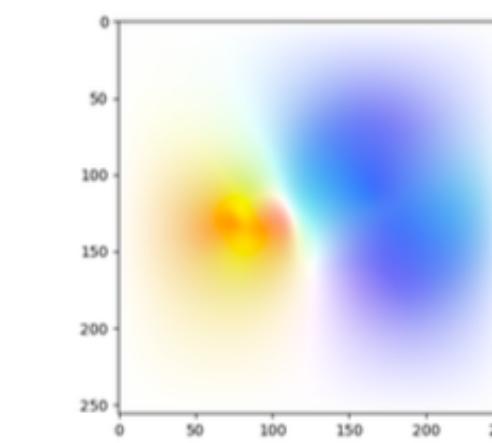
Figure 6: Second frame from synthetic test images of Gaussian's circling each other for image sizes of $2^k \times 2^k$ for $k = 6, 7, 8, 9$.



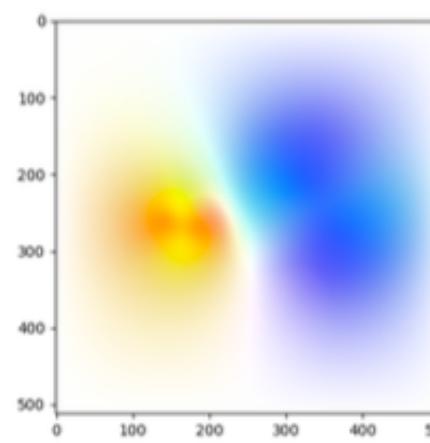
(a) $k = 6$



(b) $k = 7$



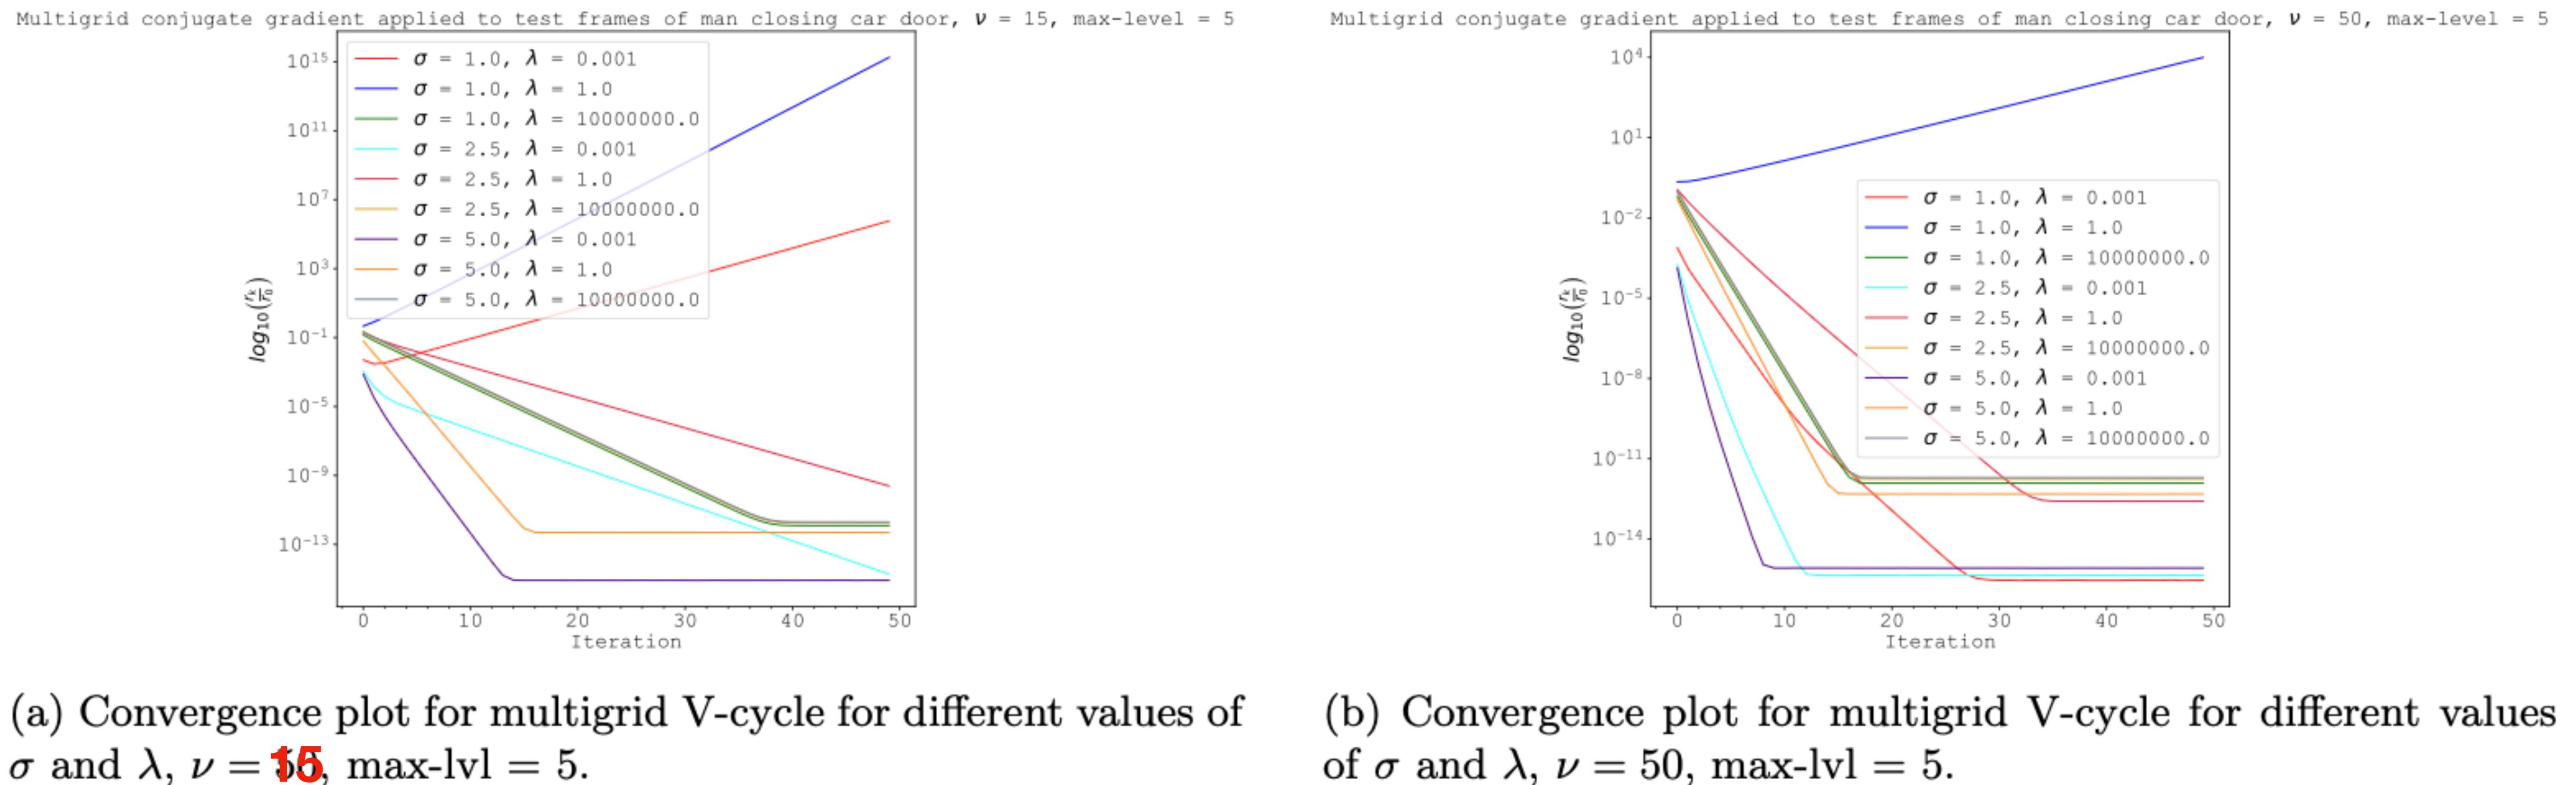
(c) $k = 8$



(d) $k = 9$

Figure 7: Resulting optical flow fields for the test case of two Gassians circling each other for image sizes of $2^k \times 2^k$ for $k = 6, 7, 8, 9$ using the sparse implementation of the conjugate gradient algorithm (on-the-grid implementation yields identical results).

Results (part IV)



(a) Convergence plot for multigrid V-cycle for different values of σ and λ , $\nu = 15$, max-lvl = 5.

(b) Convergence plot for multigrid V-cycle for different values of σ and λ , $\nu = 50$, max-lvl = 5.

Figure 4: Convergence plots for multigrid conjugate gradient for various σ and λ values.

Results (part V)

σ	λ	t	$t_{\nu=15, \text{lvl}=5}$	$t_{\nu=50, \text{lvl}=5}$	$t_{\nu=15, \text{lvl}=2}$
1.0	10^{-3}	0.0974	0.4135	1.2325	0.5367
1.0	10^0	0.0966	0.4008	1.2472	0.5177
1.0	10^7	0.0959	0.3988	1.3043	0.5087
2.5	10^{-3}	0.0961	0.3934	1.2449	0.4848
2.5	10^0	0.0969	0.3951	1.2290	0.5195
2.5	10^7	0.0977	0.3991	1.2549	0.5073
5.0	10^{-3}	0.0937	0.3891	1.2478	0.4470
5.0	10^0	0.0932	0.3826	1.3599	0.5340
5.0	10^7	0.0967	0.3982	1.2861	0.5144

Table 2: Runtimes for multigrid V-cycle for various (σ, λ) and smoothing settings.

Results (part VI)

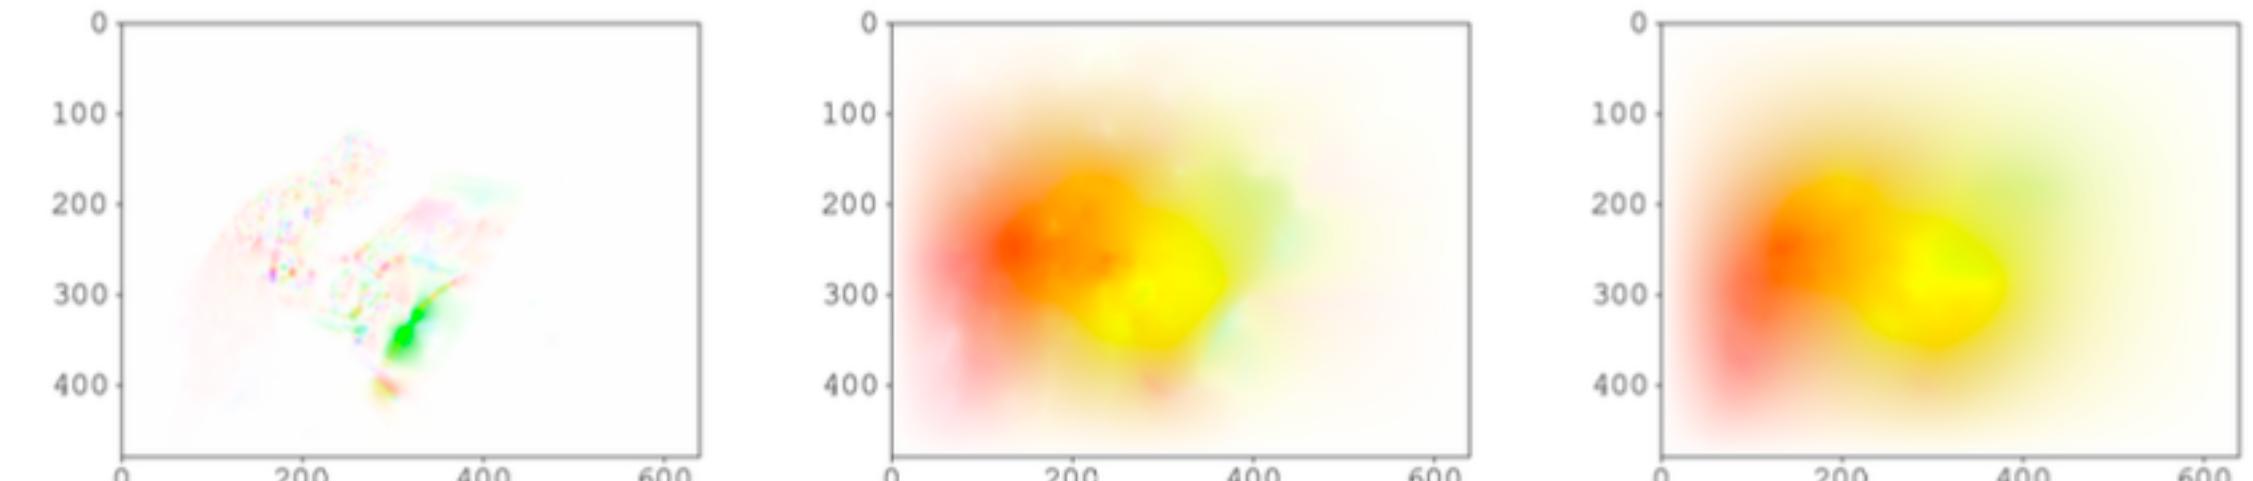
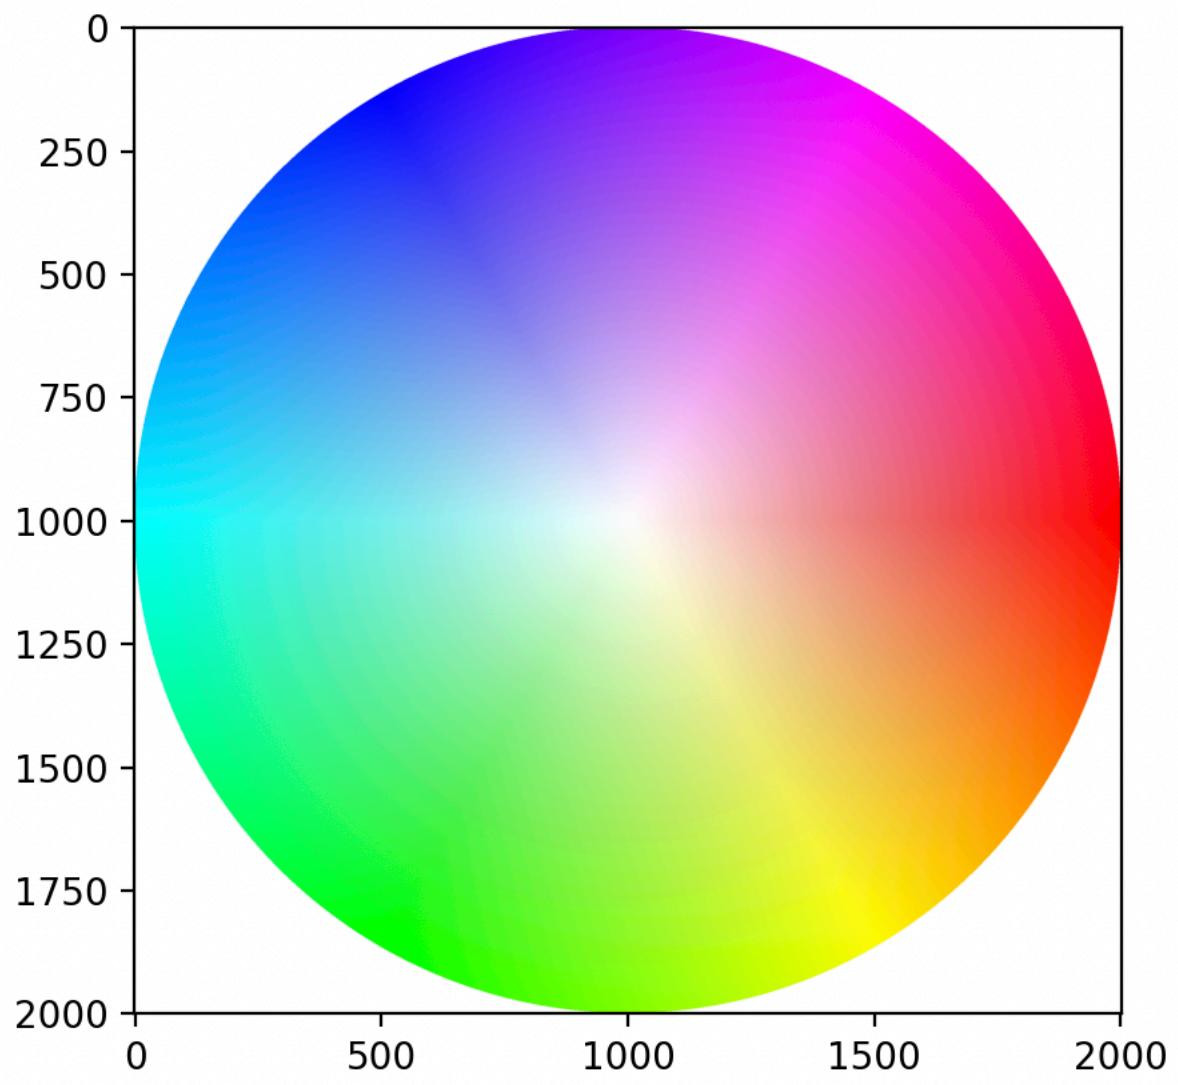


(a) First frame used for testing.



(b) Second frame used for testing.

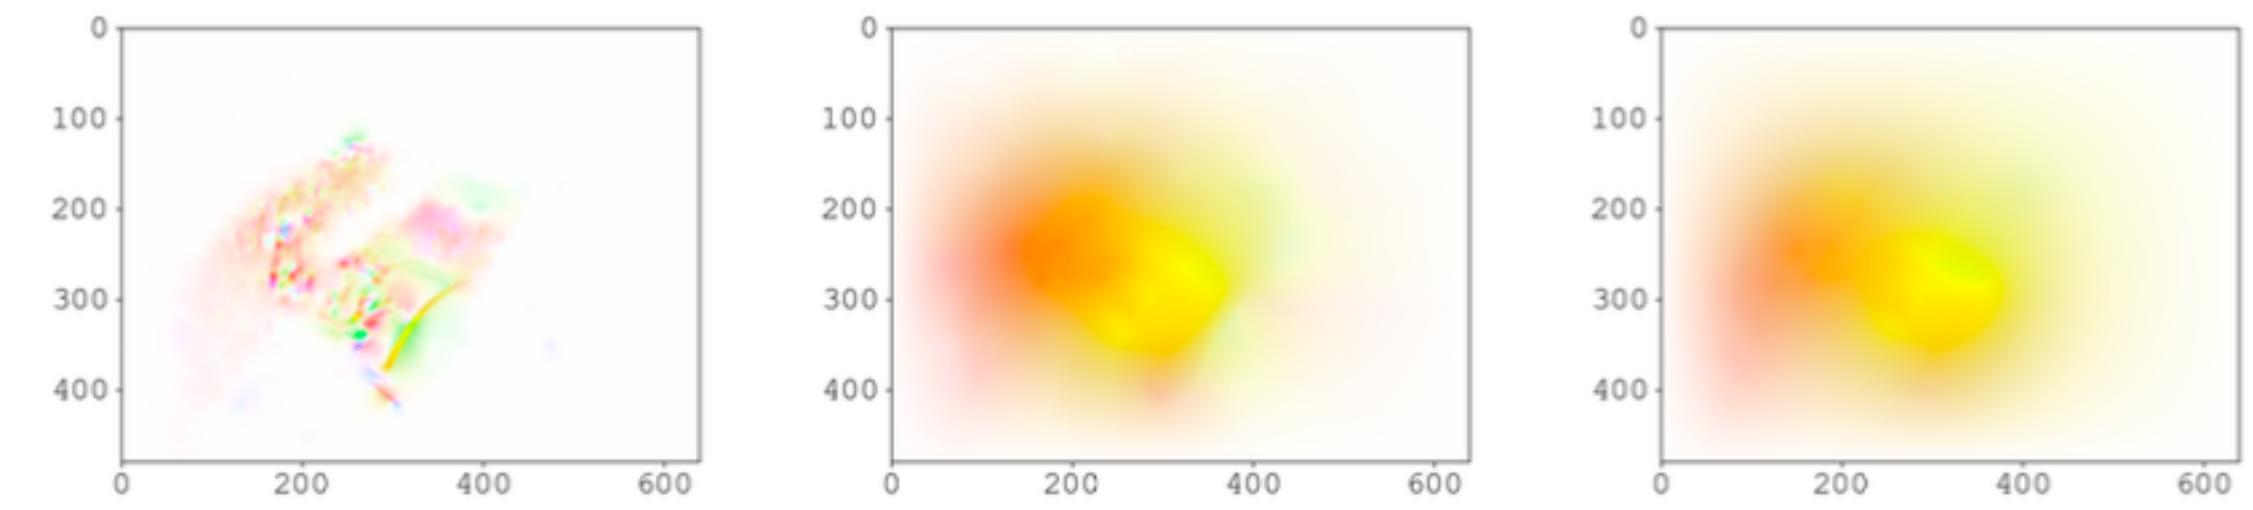
Figure 8: The two frames of a man closing his car door which were used to test the implementation.



(a) $\sigma = 1.0, \lambda = 0.001$

(b) $\sigma = 1.0, \lambda = 1.0$

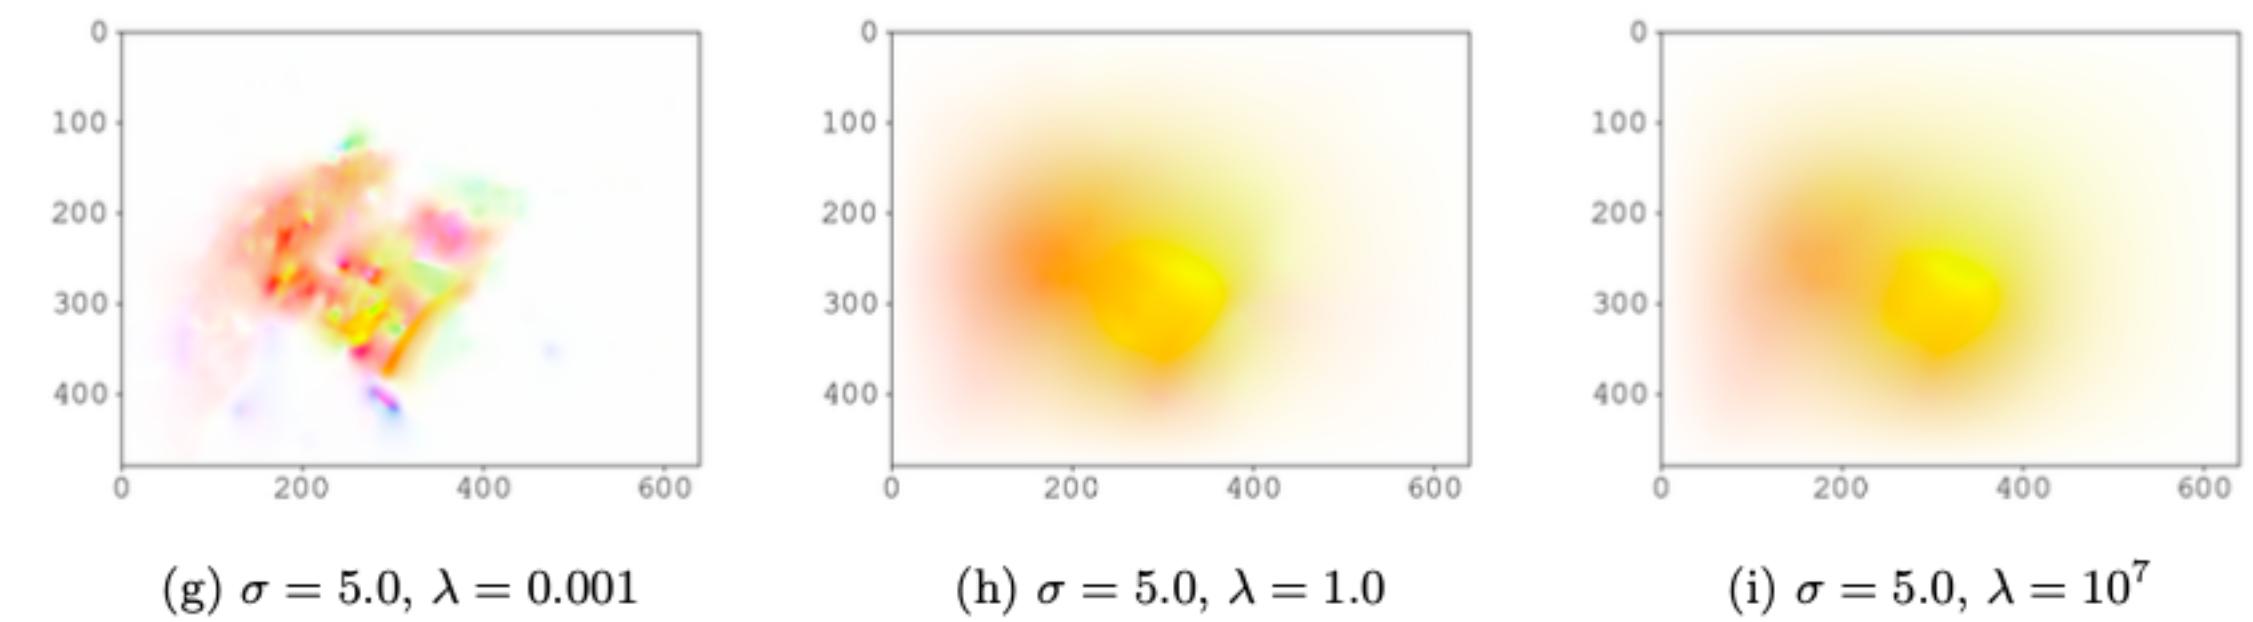
(c) $\sigma = 1.0, \lambda = 10^7$



(d) $\sigma = 2.5, \lambda = 0.001$

(e) $\sigma = 2.5, \lambda = 1.0$

(f) $\sigma = 2.5, \lambda = 10^7$



(g) $\sigma = 5.0, \lambda = 0.001$

(h) $\sigma = 5.0, \lambda = 1.0$

(i) $\sigma = 5.0, \lambda = 10^7$

Figure 9: Resulting optical flow field for the provided test frames of a man closing his car door using the multigrid conjugate gradient (on-the-grid, not sparse) algorithm. The number of iterations for the smoother was set to $\nu = 15$ and the max-level for the V-cycle was set to 5.

Results (part VII)

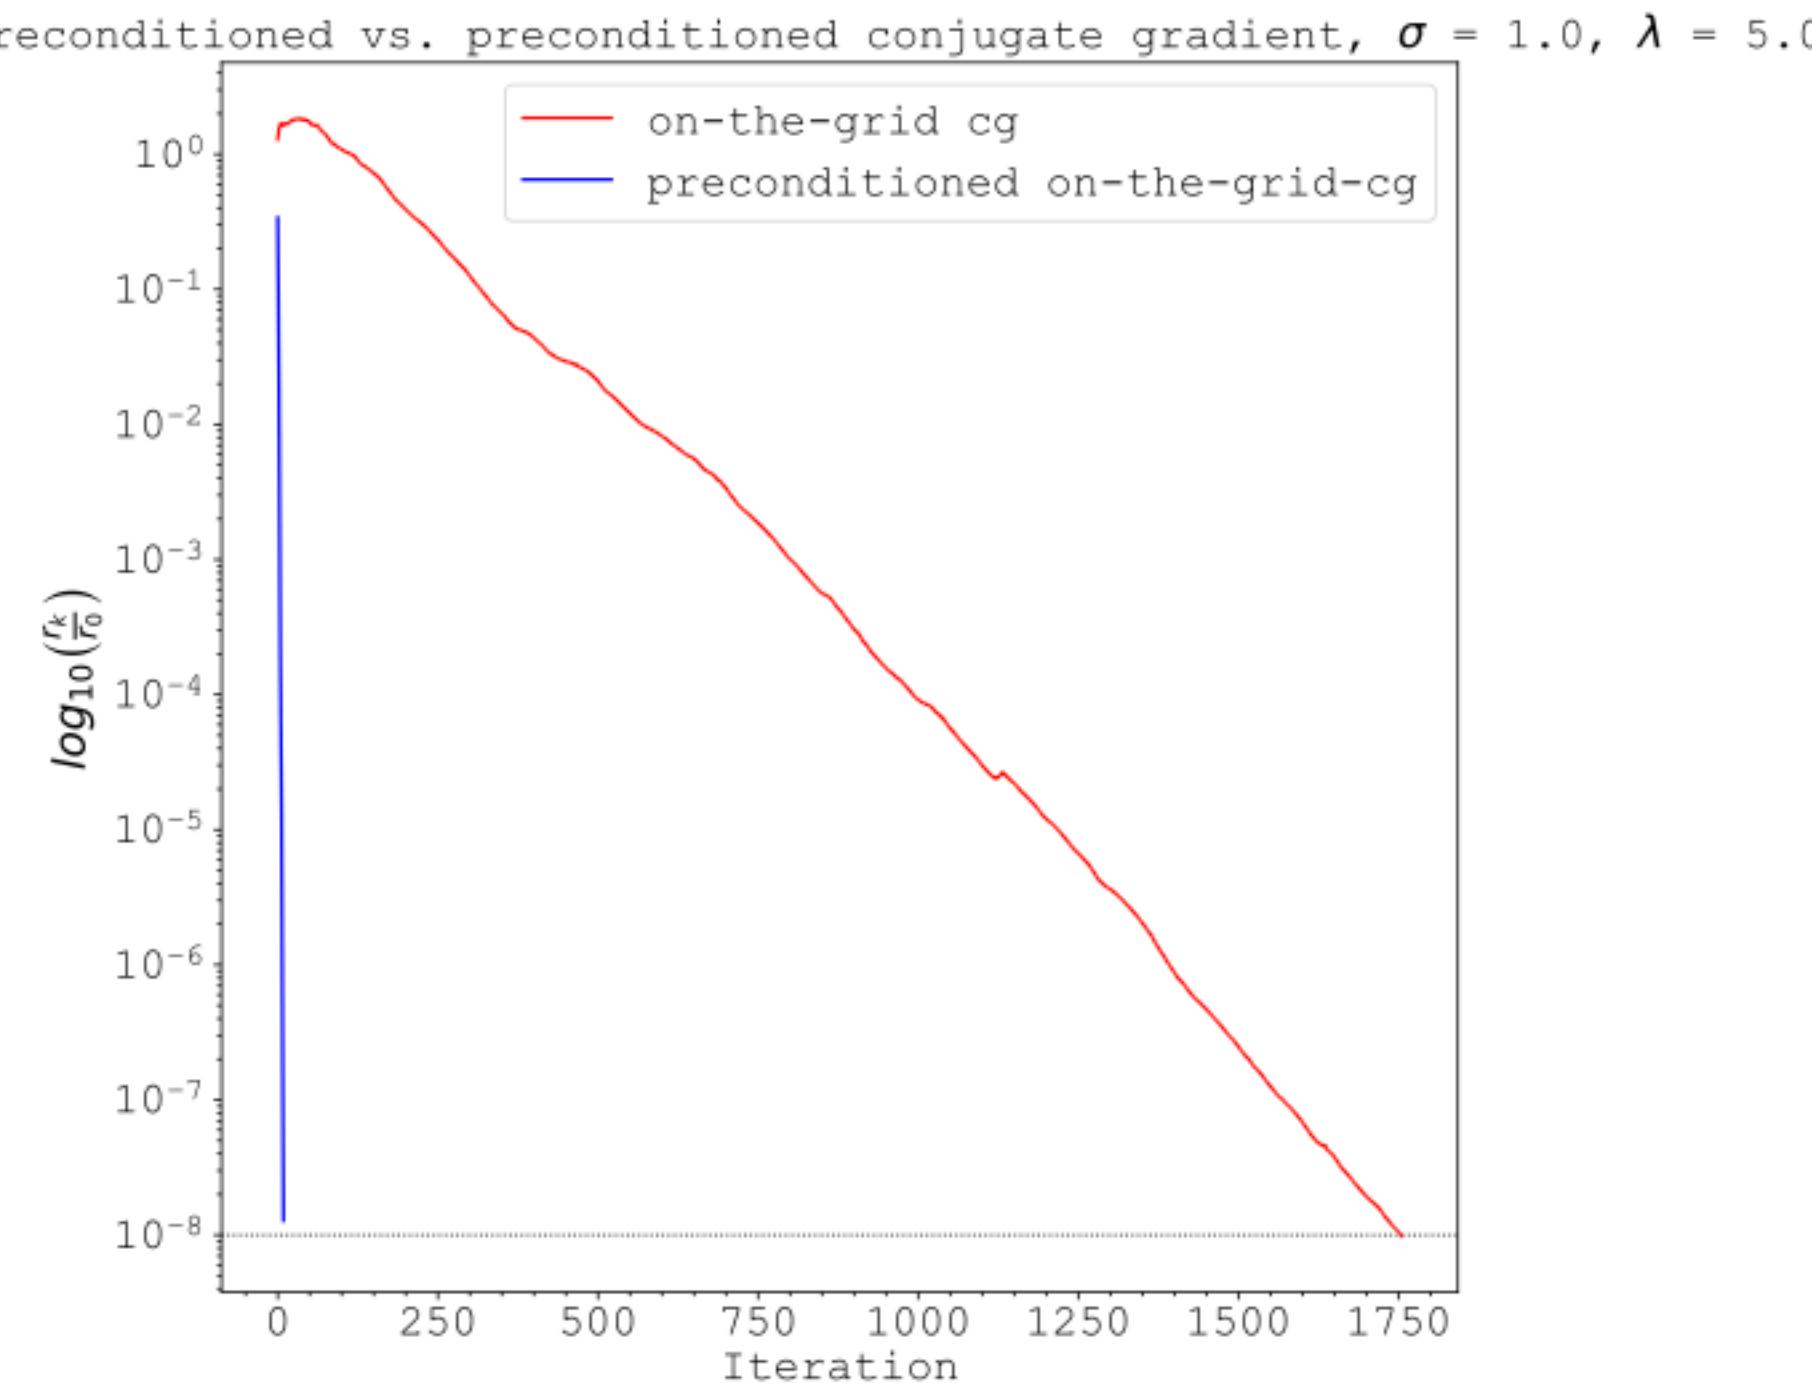


Figure 5: Convergence plot comparing the non-preconditioned on-the-grid implementation of the Conjugate Gradient algrotihm vs. the preconditioned version for the test frames of a man closing a car door.

12.210s to 4.325s.

Problems encountered

- Enforce Dirichlet B.C. ? (debugging CG padding vs work on interior)
- Update order red-black? no hahahah IxIy IyIx IxIt IyIt messup, ANNOYING BUG!!!
- Multigrid V-cycle diverges for small values of lambda (if sigma = 1, lambda < 5.0 behaves badly)
- Why h and not 2h when calling conjugate gradient on the coarse grid?

Thank you, questions?

A note on smoothing

Multigrid conjugate gradient applied to test frames of man closing car door, $\nu = 3$, max-level = 5

